

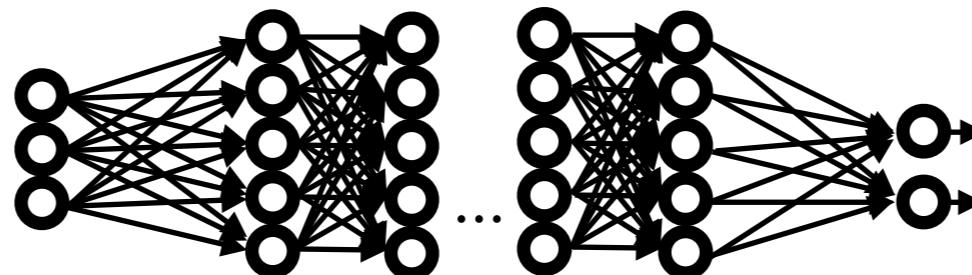
Artificial Intelligence

Taro Sekiyama

National Institute of Informatics (NII)
sekiyama@nii.ac.jp

Review

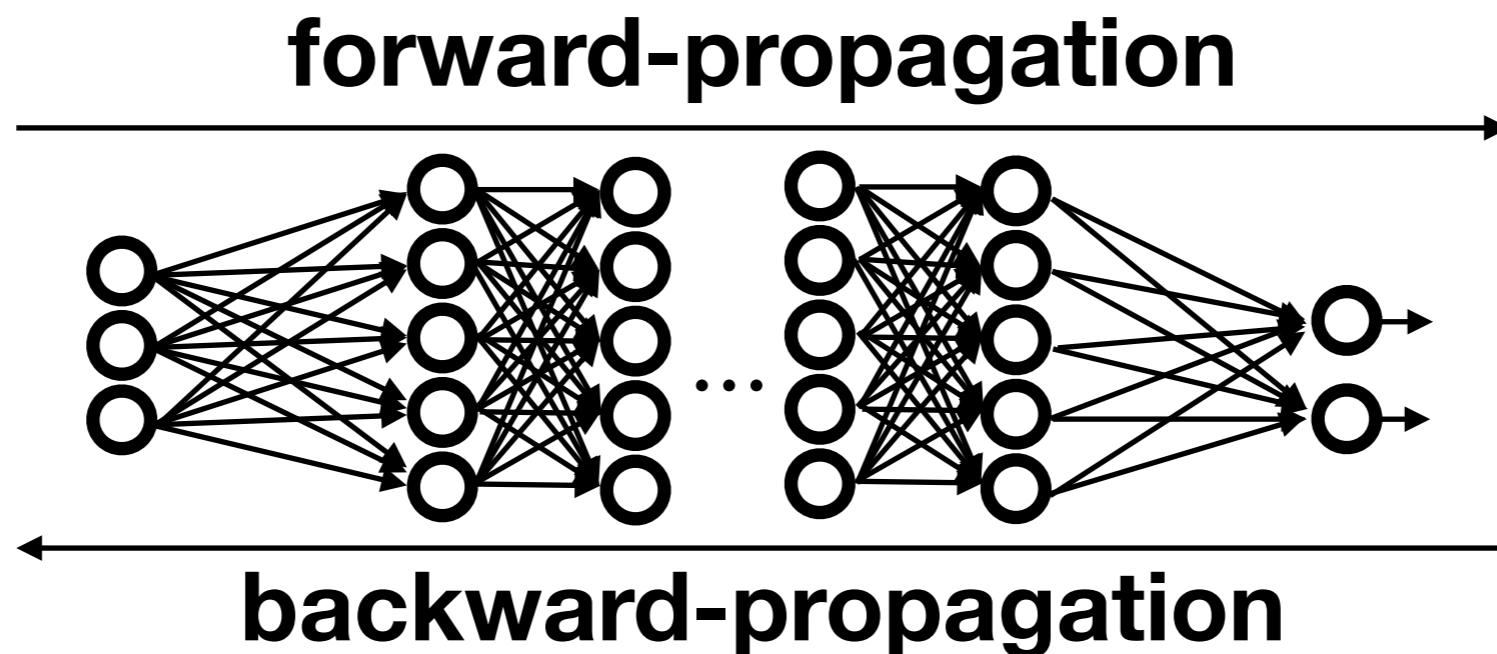
- Neural networks (NNs) are powerful, but difficult-to-amend, ML models
- Simple feedforward NN model



- Including multiple hidden layers
(called ***fully connected***)
- Neuron: basic computational unit
 - Linear transformation + activation function

Review

- Training of NNs is to optimize parameters of neurons
 - The gradient descent algorithm to optimize is called *back-propagation*



Review

- The vanishing gradient problem
 - Gradients in deeper NNs tend to approach to zero
- Important to strike a balance between expressivity and trainability

Challenges of NNs

1. Hard to find optimal parameters even by gradient descent
 - The expressivity of FNNs makes their cost functions non-convex
 - Important to design ***features*** that makes optimization easier
2. How do we handle sequential data without the vanishing gradient problem?

This lecture

Introducing two special types of NNs

- Convolutional neural networks (CNNs)
 - Specialized in image processing
 - Good at extracting features from images
- Recurrent neural networks (RNNs)
 - Specialized in sequential data
 - Relaxing the vanishing gradient problem

Success of CNNs

■ Image recognition



Images from: <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

Success of CNNs

■ Image recognition



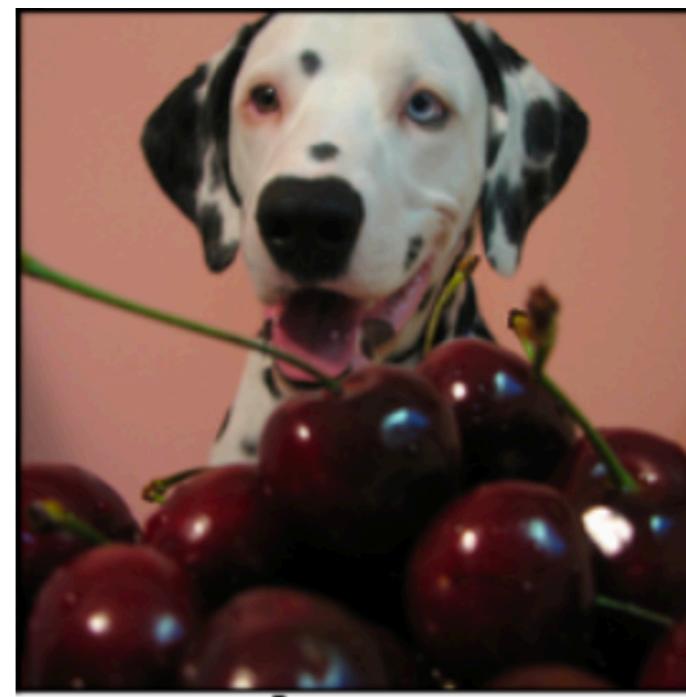
mite

container ship



motor scooter

leopard

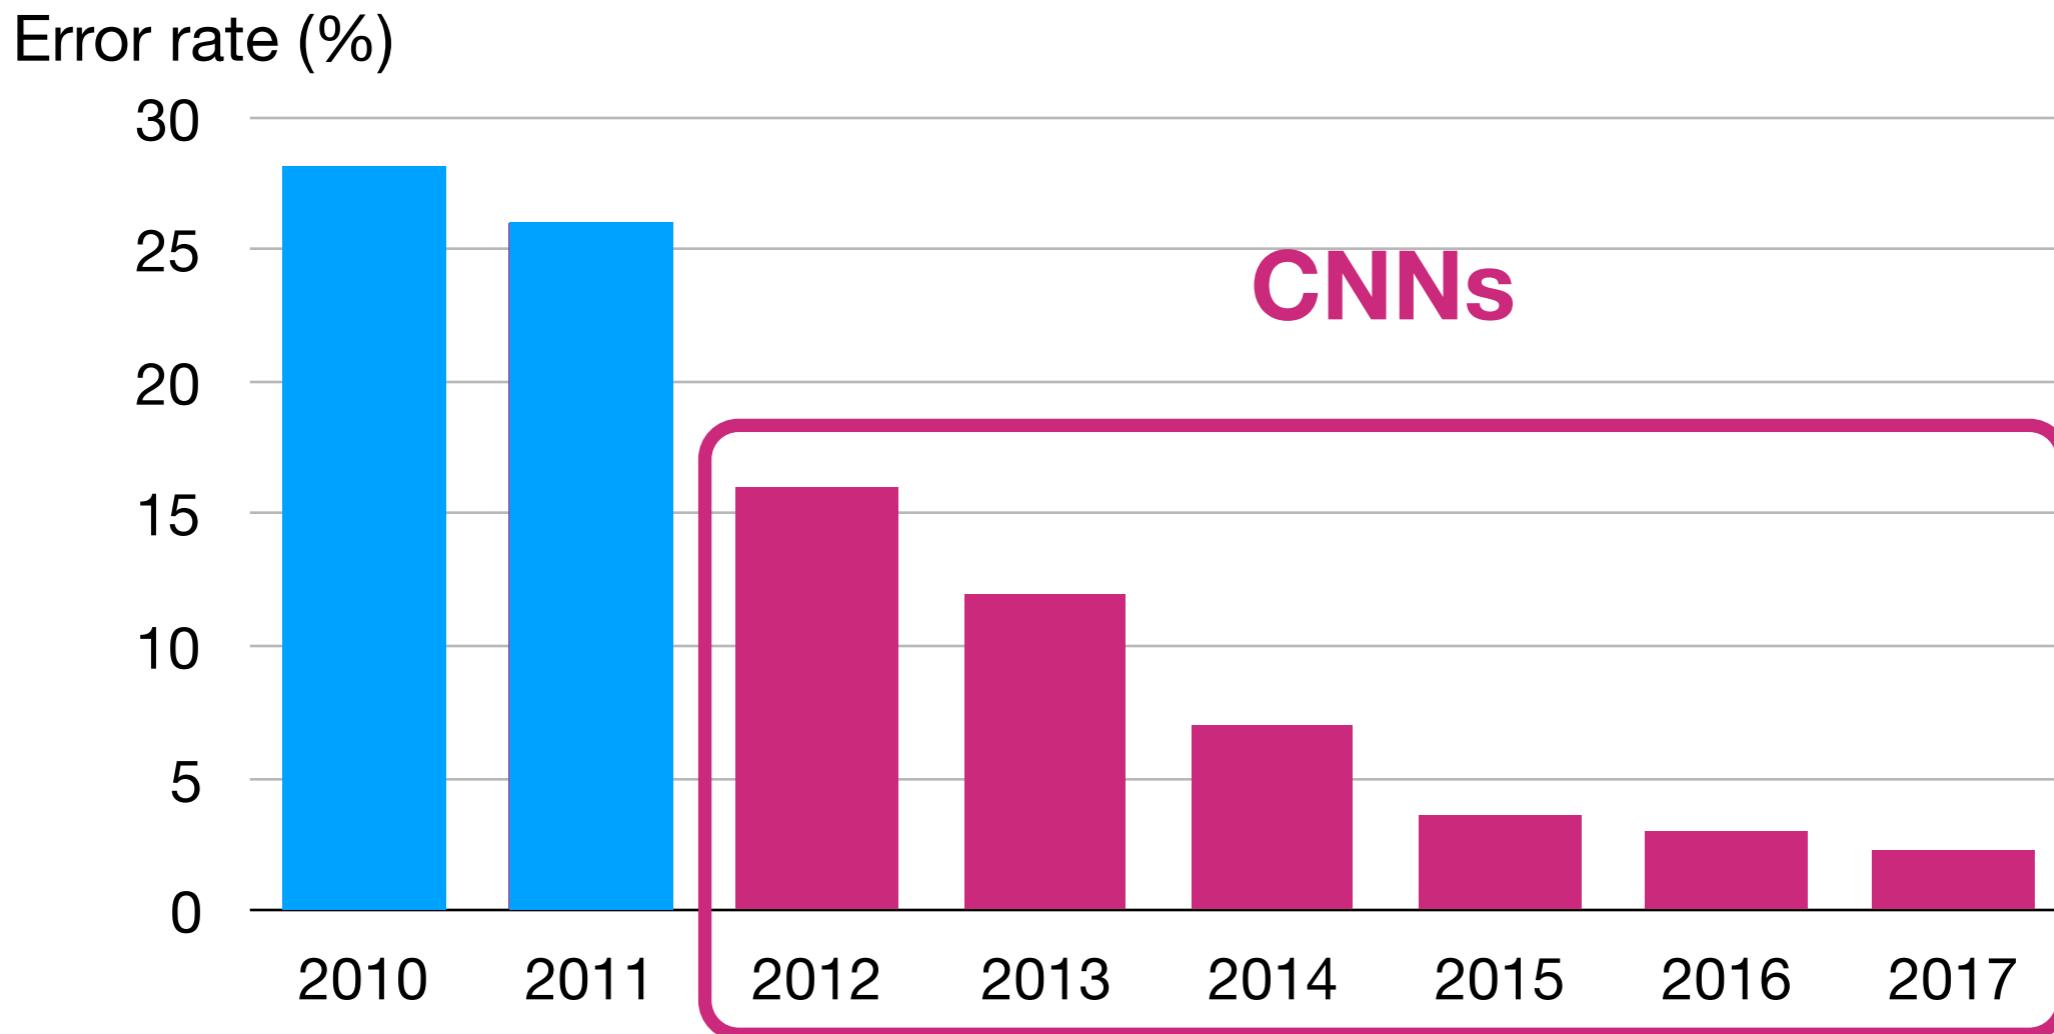


cherry

Images from: <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

Success of CNNs

- CNNs have been winning in ILSVRC (ImageNet Large Scale Visual Recognition Competition) since 2012



Success of CNNs

- Object detection

<https://towardsdatascience.com/object-detection-with-10-lines-of-code-d6cb4d86f606>

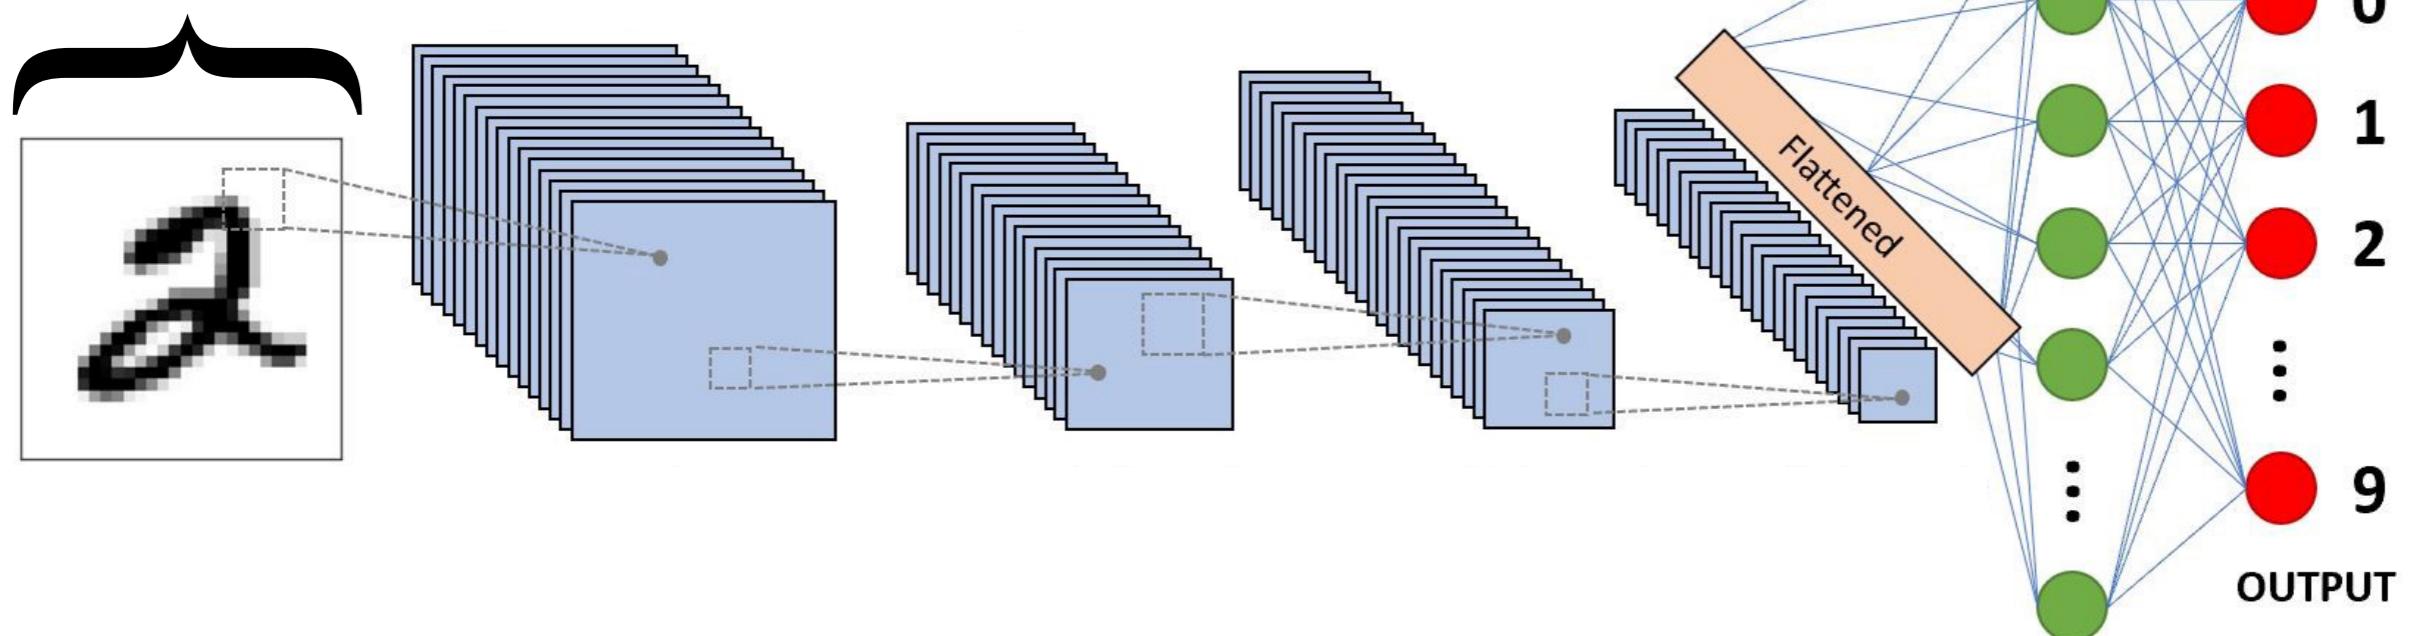
Success of CNNs

- Image generation/transformation

[https://twitter.com/pixivsketch/status/
890547530810875904/photo/1](https://twitter.com/pixivsketch/status/890547530810875904/photo/1)

Convolution neural networks

Input image
 $\in \mathbb{R}^{H \times W \times C}$

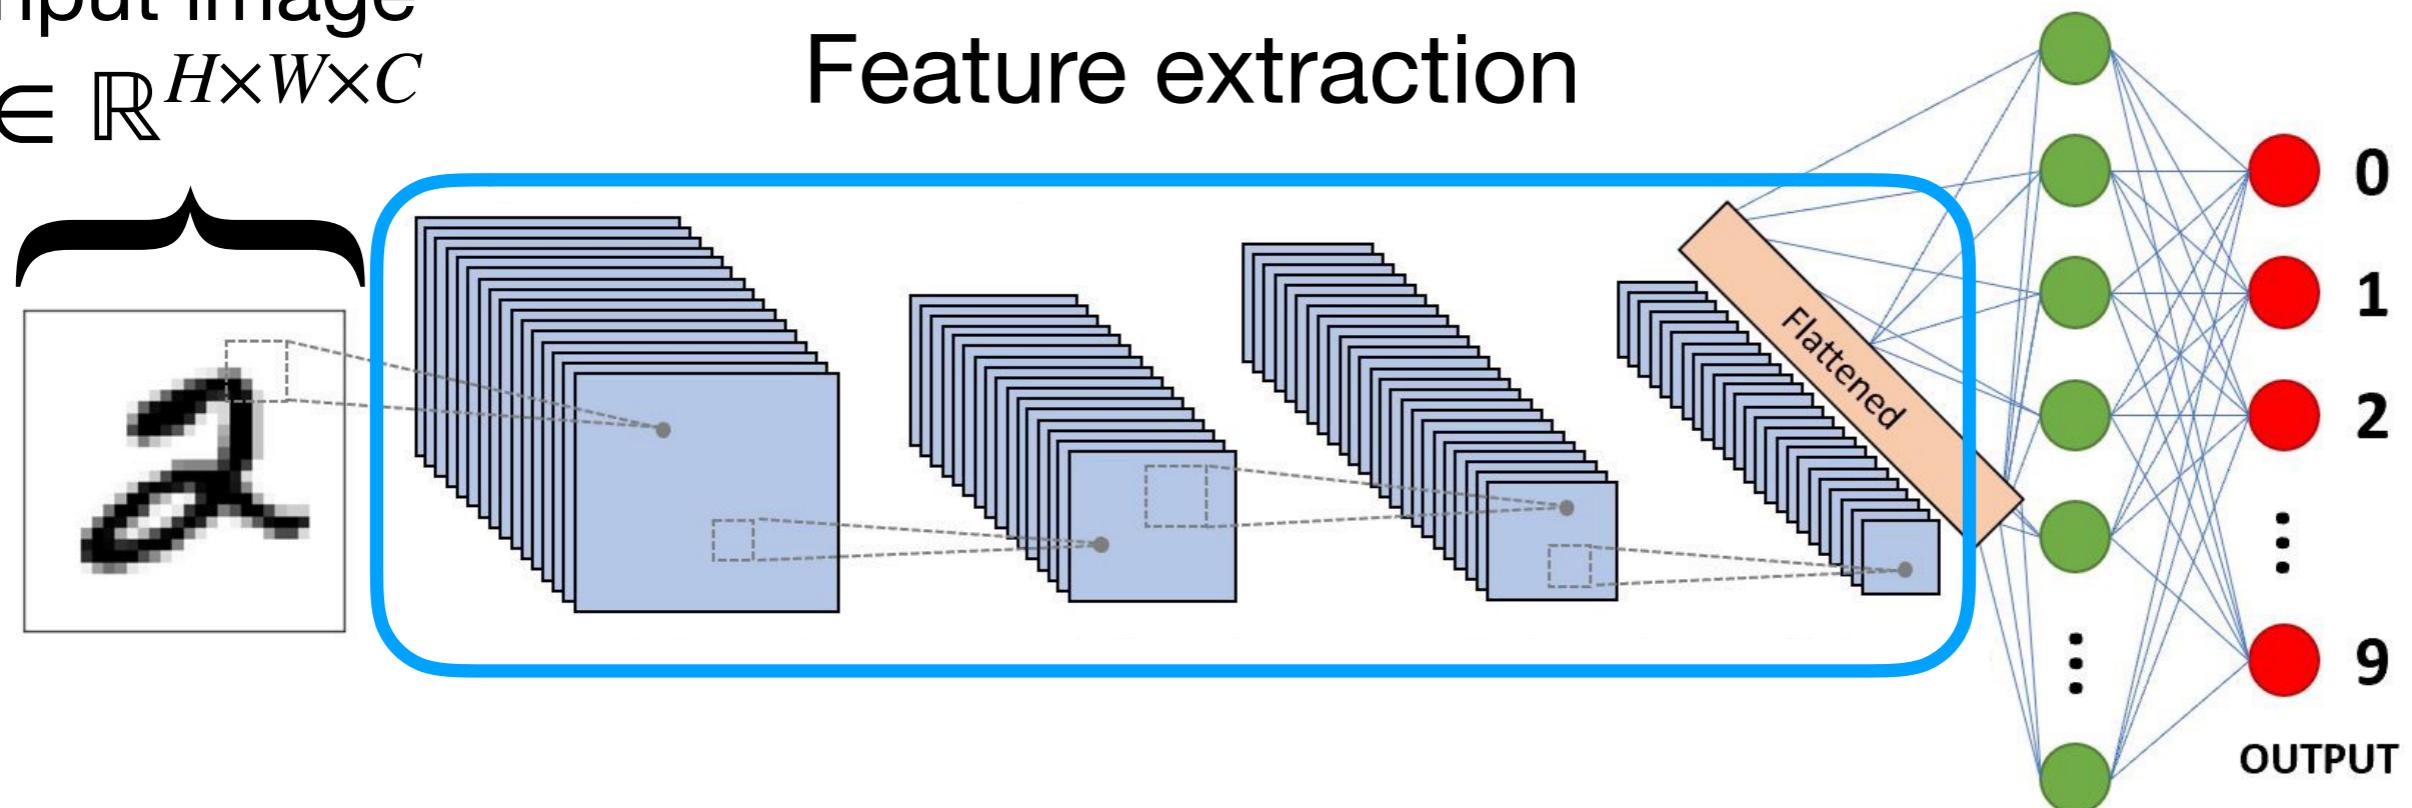


The image from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Convolution neural networks

Input image
 $\in \mathbb{R}^{H \times W \times C}$

Feature extraction



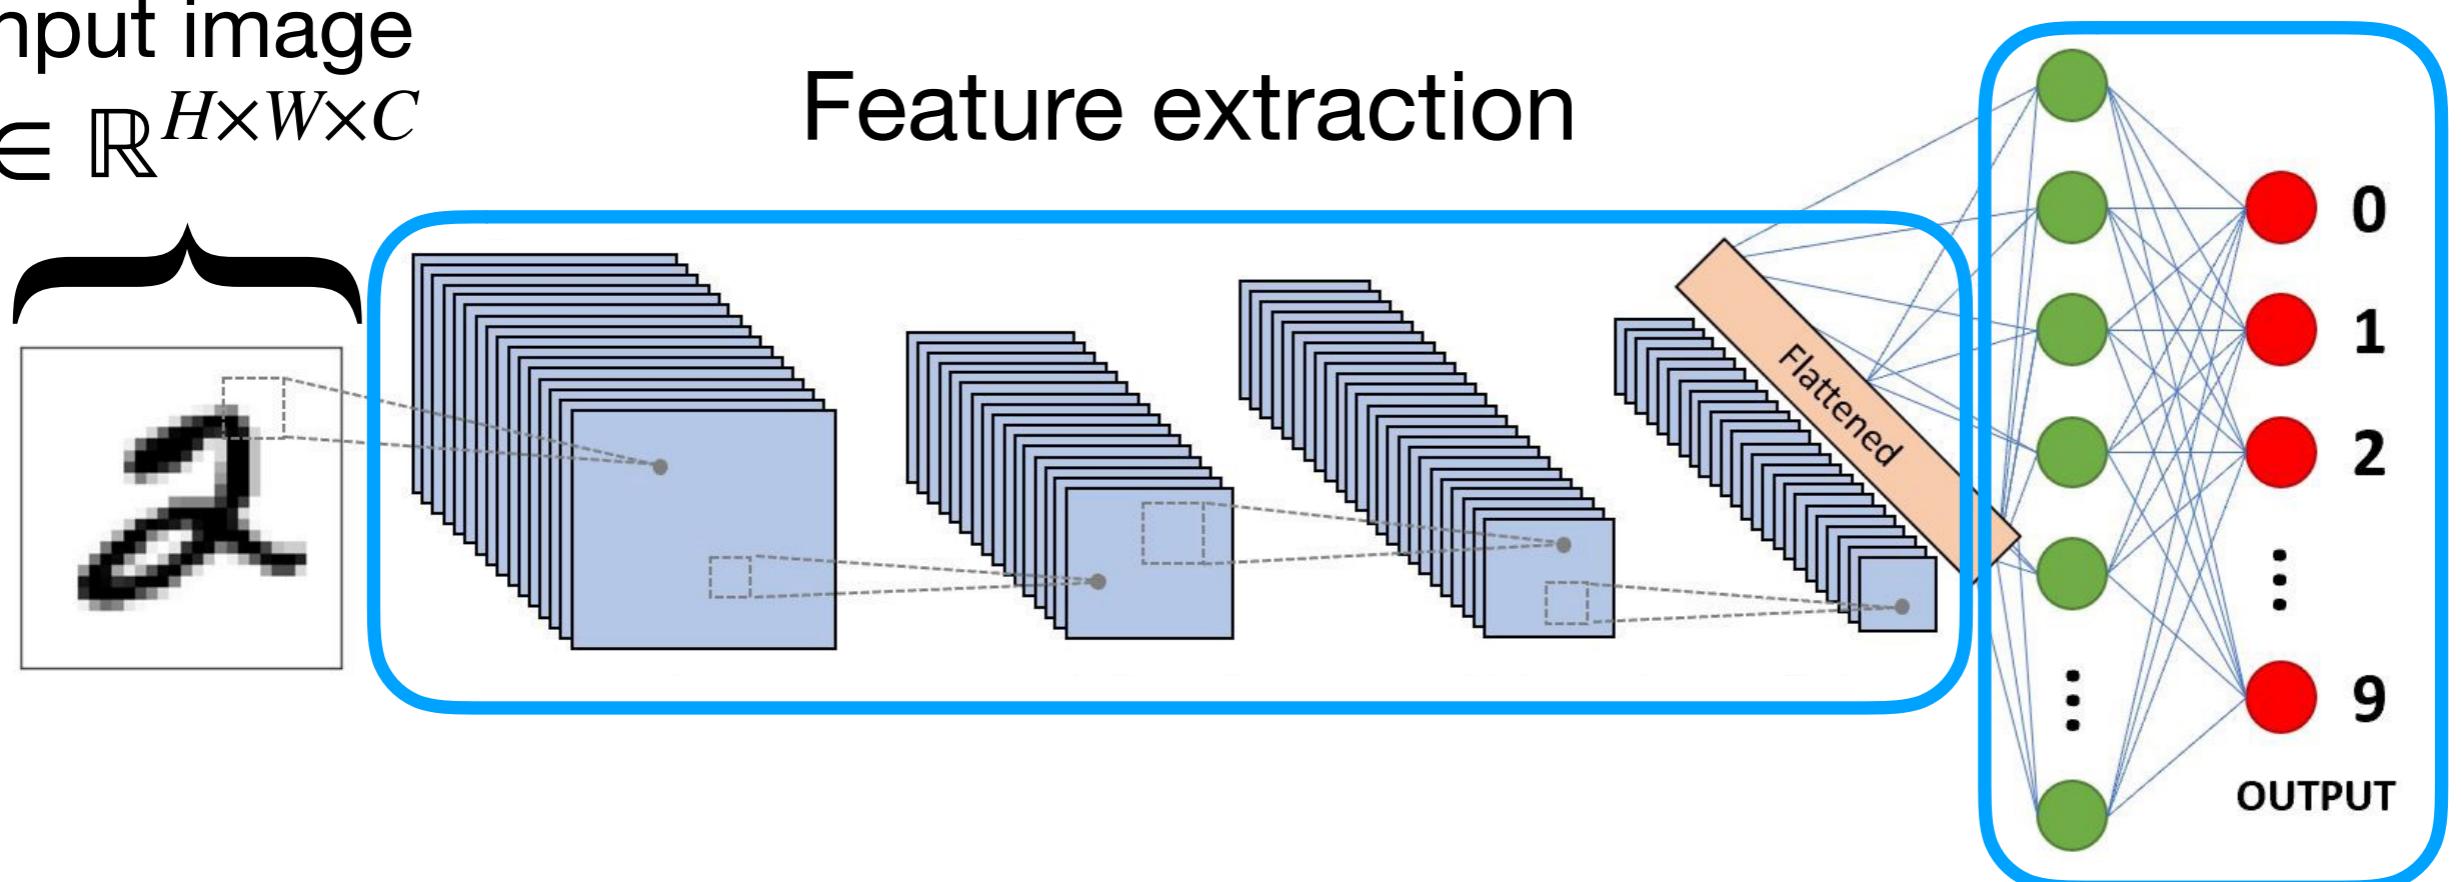
The image from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Convolution neural networks

Input image
 $\in \mathbb{R}^{H \times W \times C}$

Feature extraction

Classification
by fully connected layers
+ extracted features



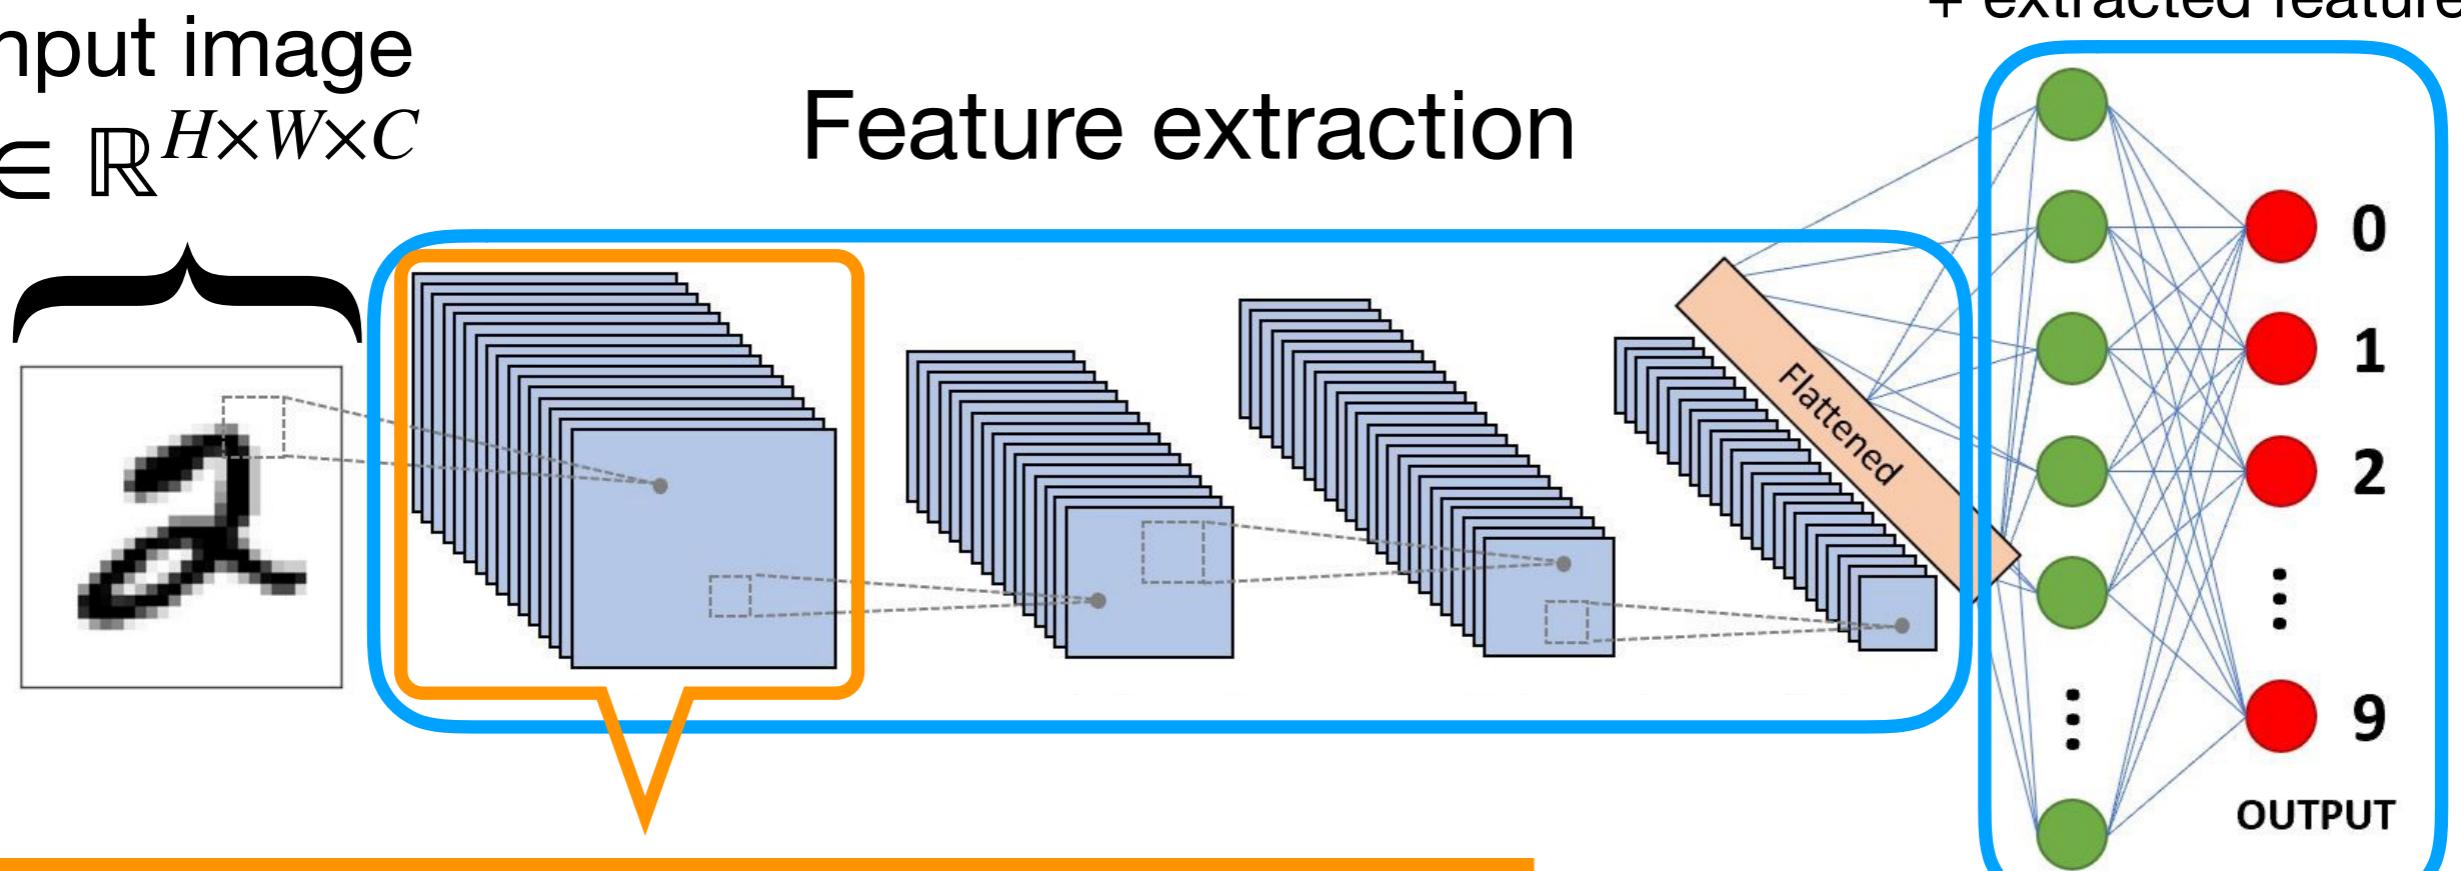
The image from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Convolution neural networks

Input image
 $\in \mathbb{R}^{H \times W \times C}$

Feature extraction

Classification
by fully connected layers
+ extracted features

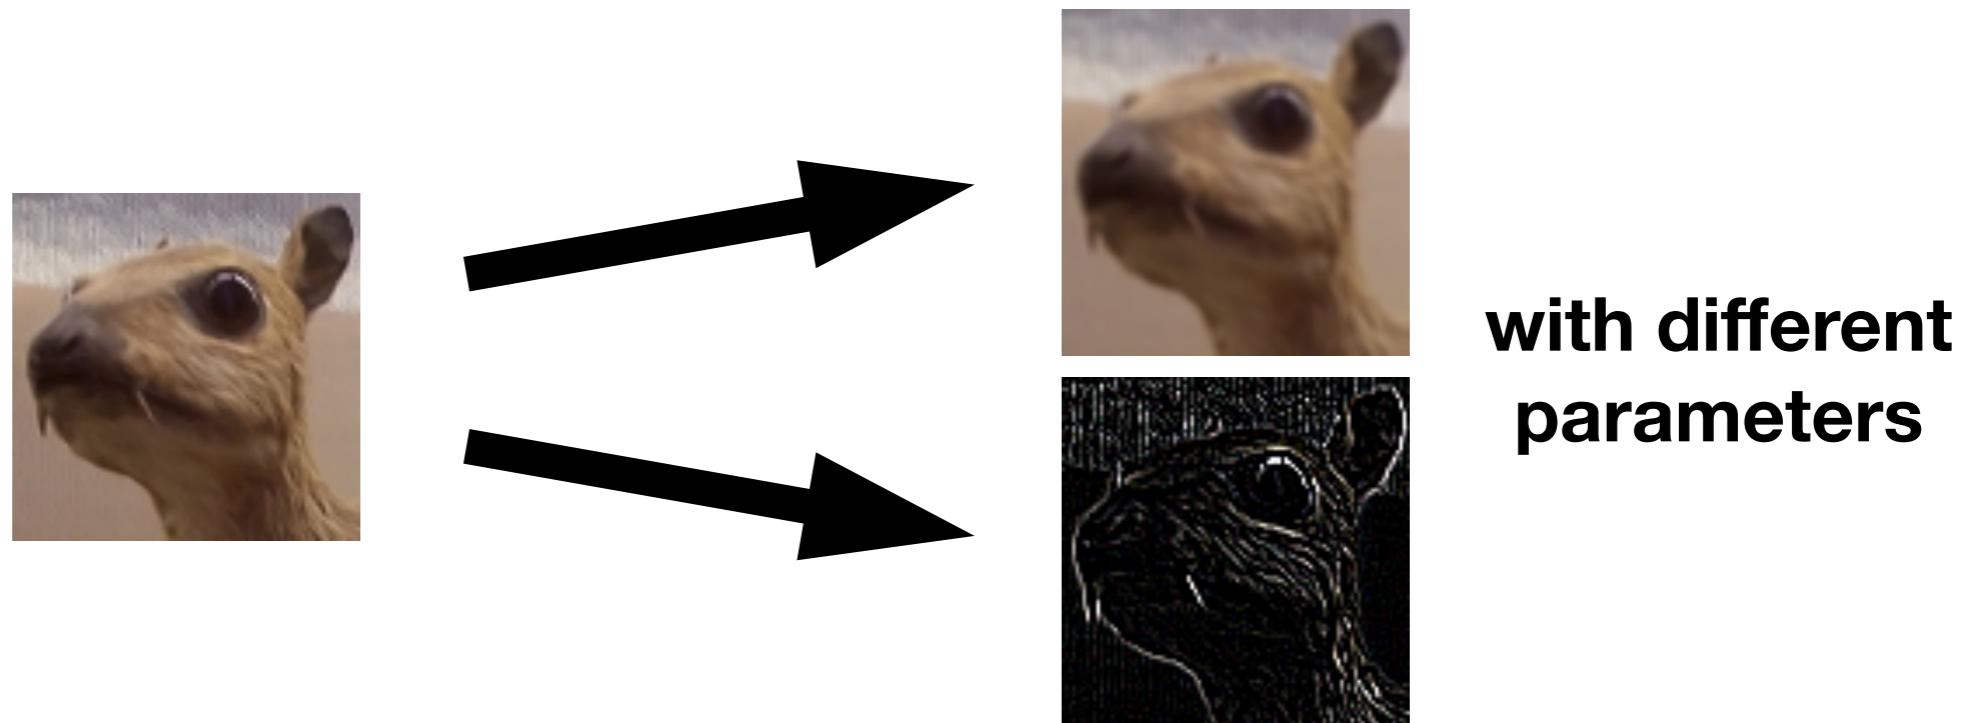


Features extracted
by a convolutional + pooling layer

The image from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Convolution

- Image transformation by filtering

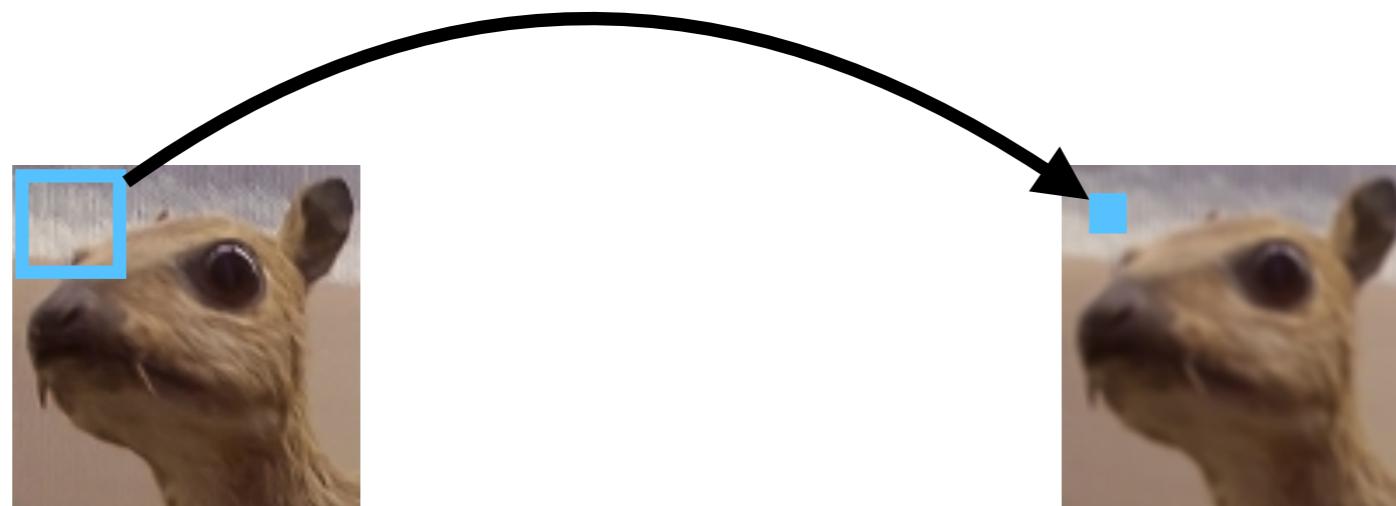


- Convolutional layers learns transformation parameters that produce image features useful for classification

Images from: [https://en.wikipedia.org/wiki/Kernel_\(image_processing\)](https://en.wikipedia.org/wiki/Kernel_(image_processing))

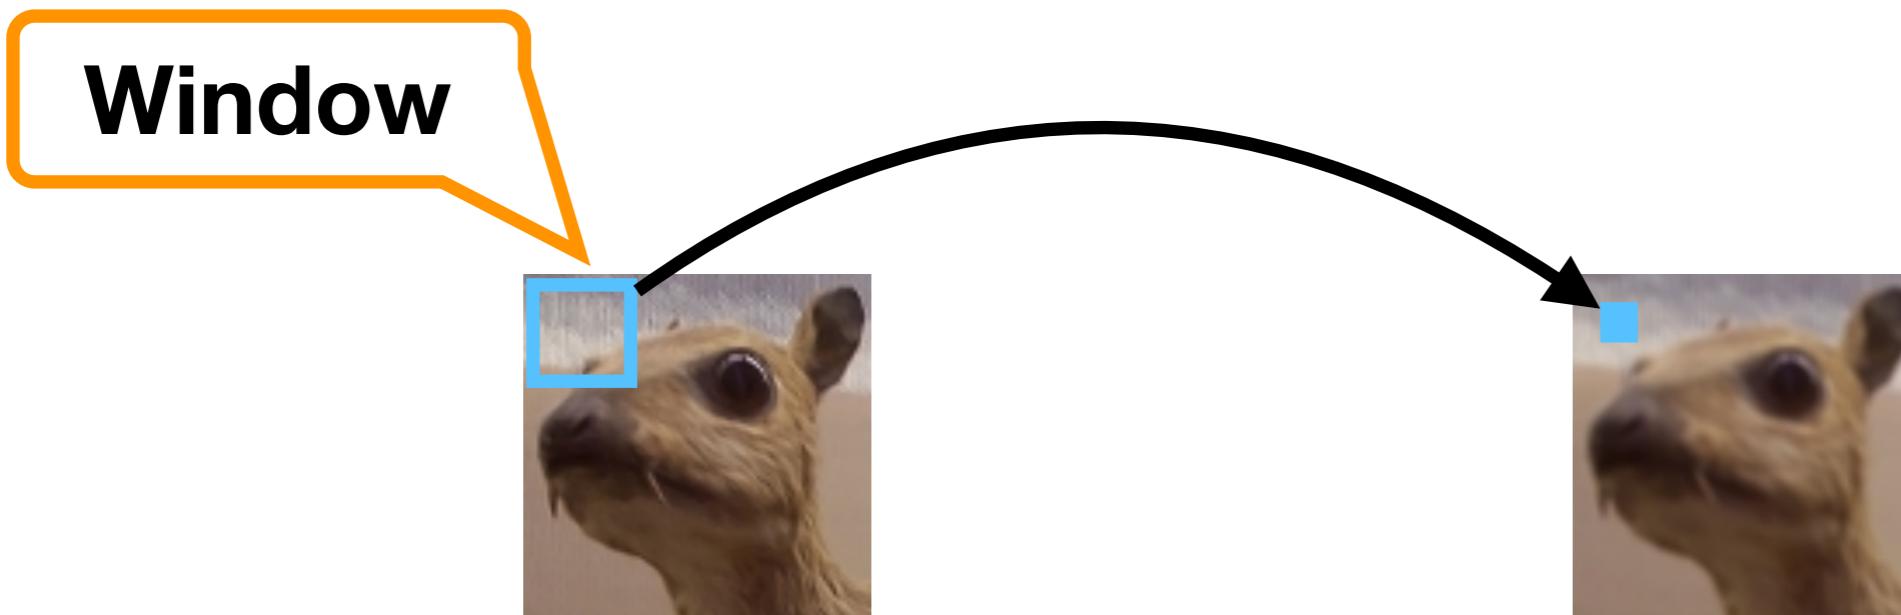
Convolution

- Each pixel of the output image is computed with a small region of the input image



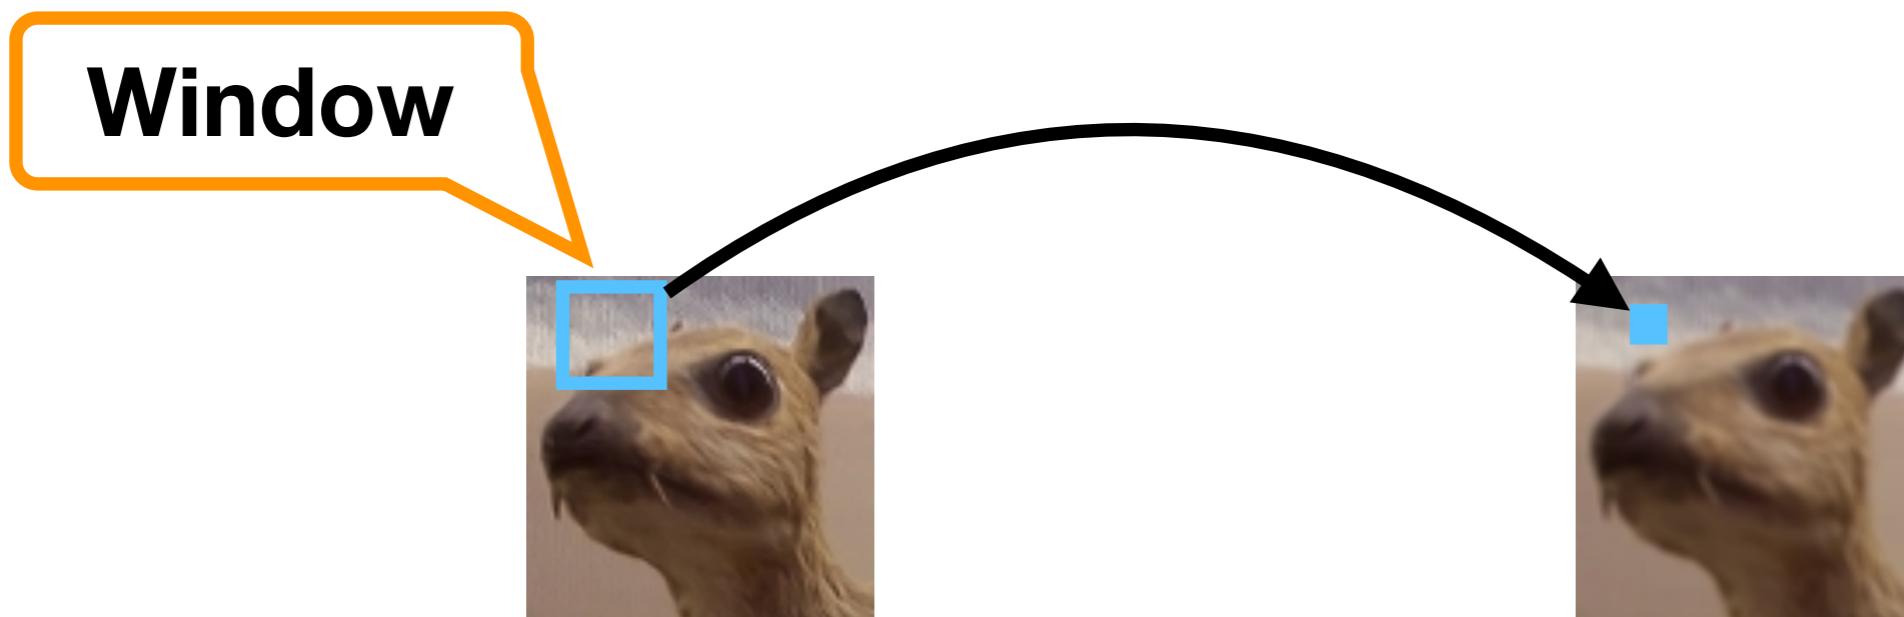
Convolution

- Each pixel of the output image is computed with a small region of the input image



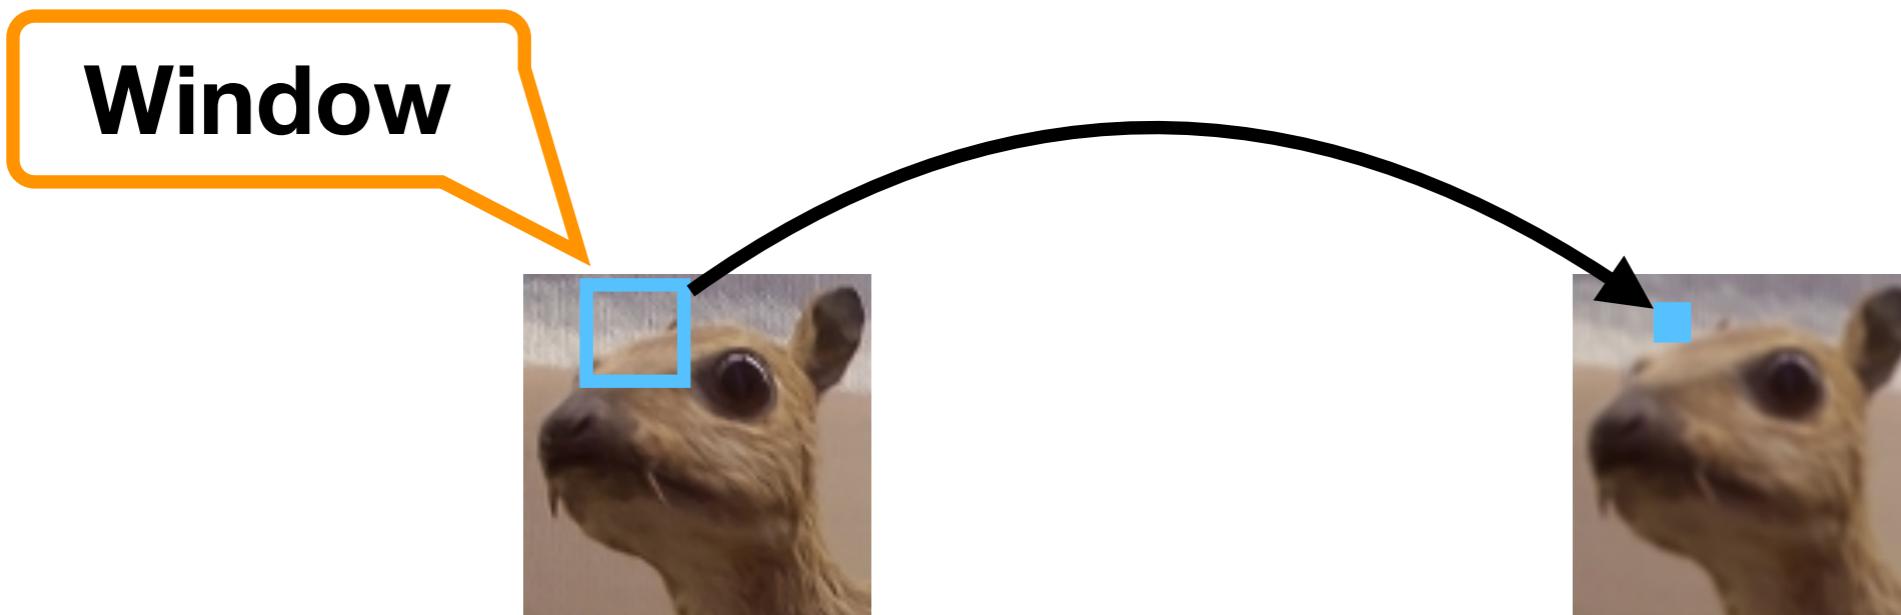
Convolution

- Each pixel of the output image is computed with a small region of the input image



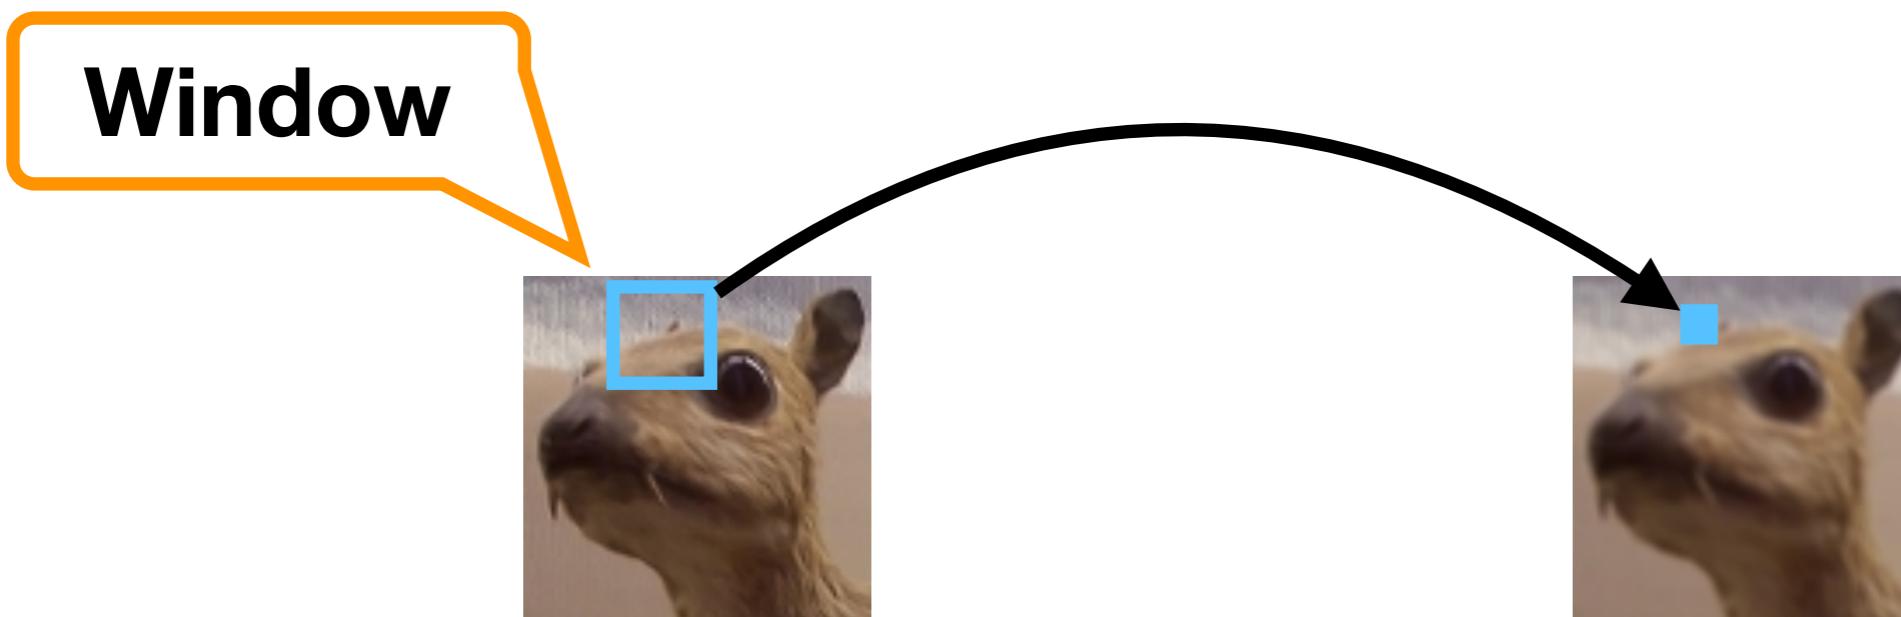
Convolution

- Each pixel of the output image is computed with a small region of the input image



Convolution

- Each pixel of the output image is computed with a small region of the input image



Running example

Window

$$W = \begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$
$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix} \xrightarrow{\hspace{1cm}} \begin{bmatrix} 12 & 8 \\ 35 & 42 \end{bmatrix}$$

Running example

Window

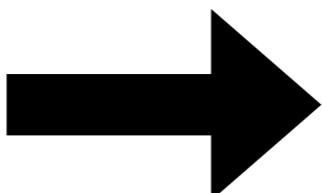
$$W = \begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$
$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix} \xrightarrow{\hspace{1cm}} \begin{bmatrix} 12 & 8 \\ 35 & 42 \end{bmatrix}$$

Running example

$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix}$$

Window

$$W = \begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$



The sum of

values of

$$\begin{bmatrix} 1 & -1 \\ 5 & 0 \end{bmatrix}$$

weighted by

$$\begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$

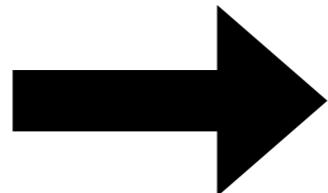
$$\begin{bmatrix} 12 & 8 \\ 35 & 42 \end{bmatrix}$$

Running example

$$\begin{bmatrix} 1 & \boxed{-1 & 3} \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix}$$

Window

$$W = \begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$



The sum of

values of

$$\begin{bmatrix} -1 & 3 \\ 0 & 2 \end{bmatrix}$$

weighted by

$$\begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$

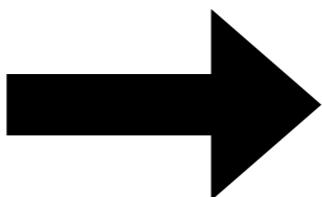
$$\begin{bmatrix} 12 & \boxed{8} \\ 35 & 42 \end{bmatrix}$$

Running example

$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix}$$

Window

$$W = \begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$



The sum of

values of

$$\begin{bmatrix} 5 & 0 \\ -2 & 6 \end{bmatrix}$$

weighted by

$$\begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$

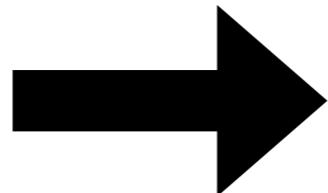
$$\begin{bmatrix} 12 & 8 \\ 35 & 42 \end{bmatrix}$$

Running example

$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix}$$

Window

$$W = \begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$



The sum of

values of

$$\begin{bmatrix} 0 & 2 \\ 6 & 7 \end{bmatrix}$$

weighted by

$$\begin{bmatrix} 3 & 1 \\ 2 & 4 \end{bmatrix}$$

$$\begin{bmatrix} 12 & 8 \\ 35 & 42 \end{bmatrix}$$

Example

Window

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \longrightarrow$$



Convolution

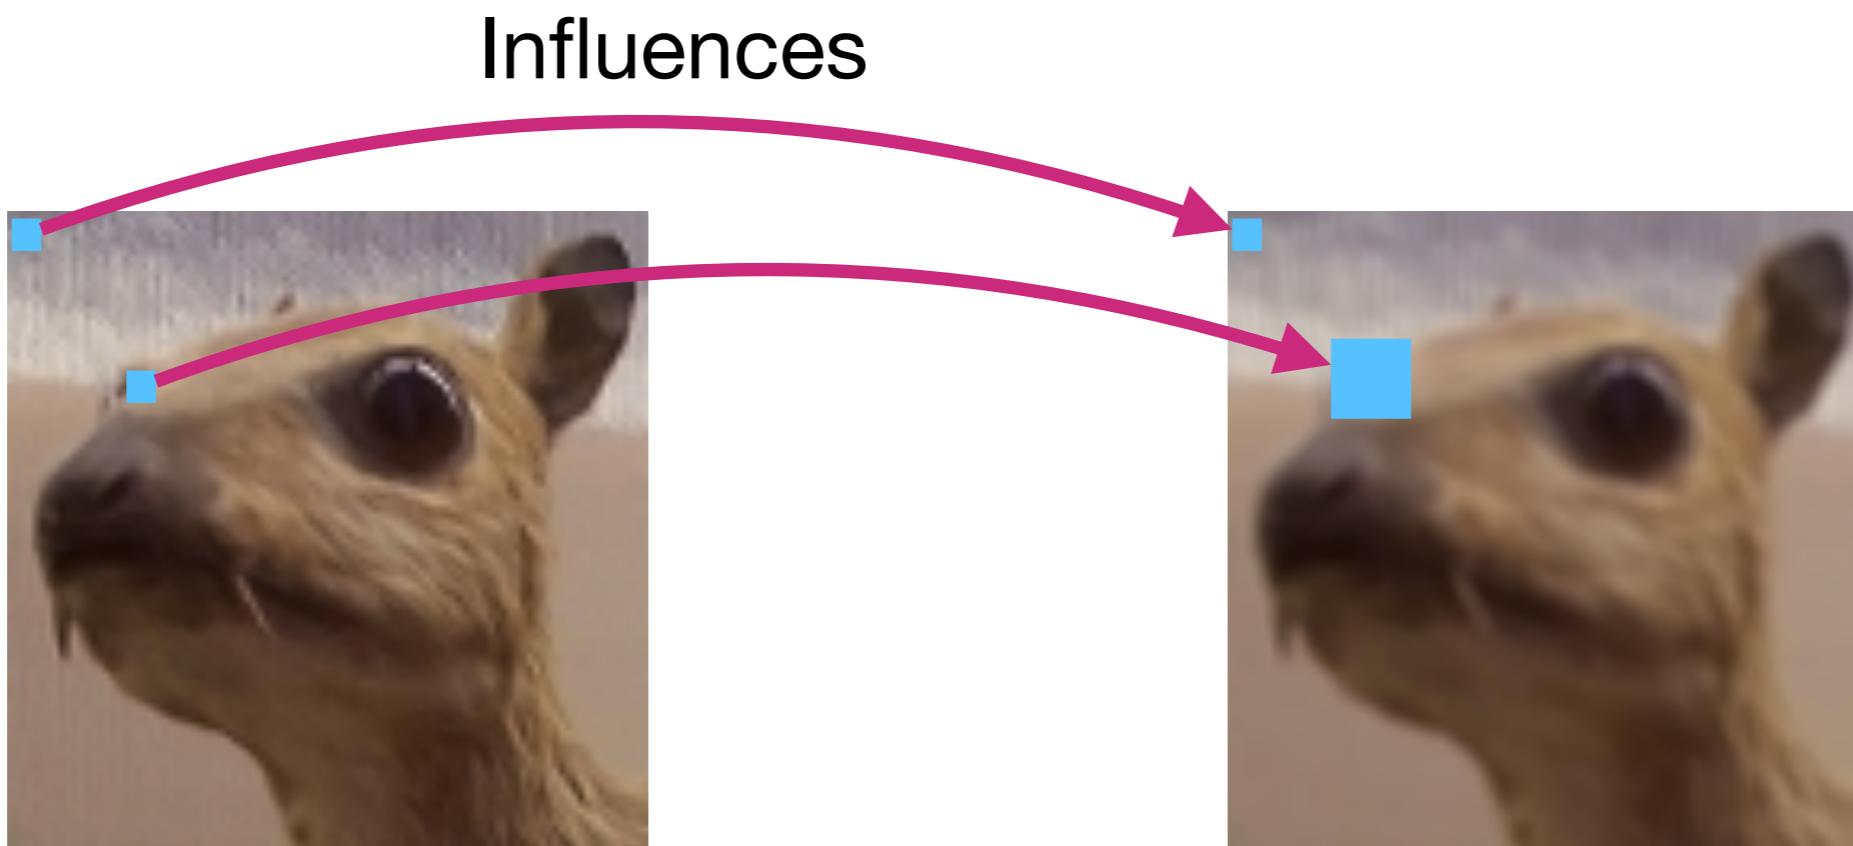
■ Input

- Features $I : H \times W \times C$ array
 - H : image height, W : image width, C : the # of features
 - Window $K : h \times w \times c$ array
- ## ■ Output
- $(H - h + 1) \times (W - w + 1) \times (C - c + 1)$ array A s.t.

$$A_{i,j,k} = \sum_{i'=1}^h \sum_{j'=1}^w \sum_{k'=1}^c I_{i+i', j+j', k+k'} K_{i', j', k'}$$

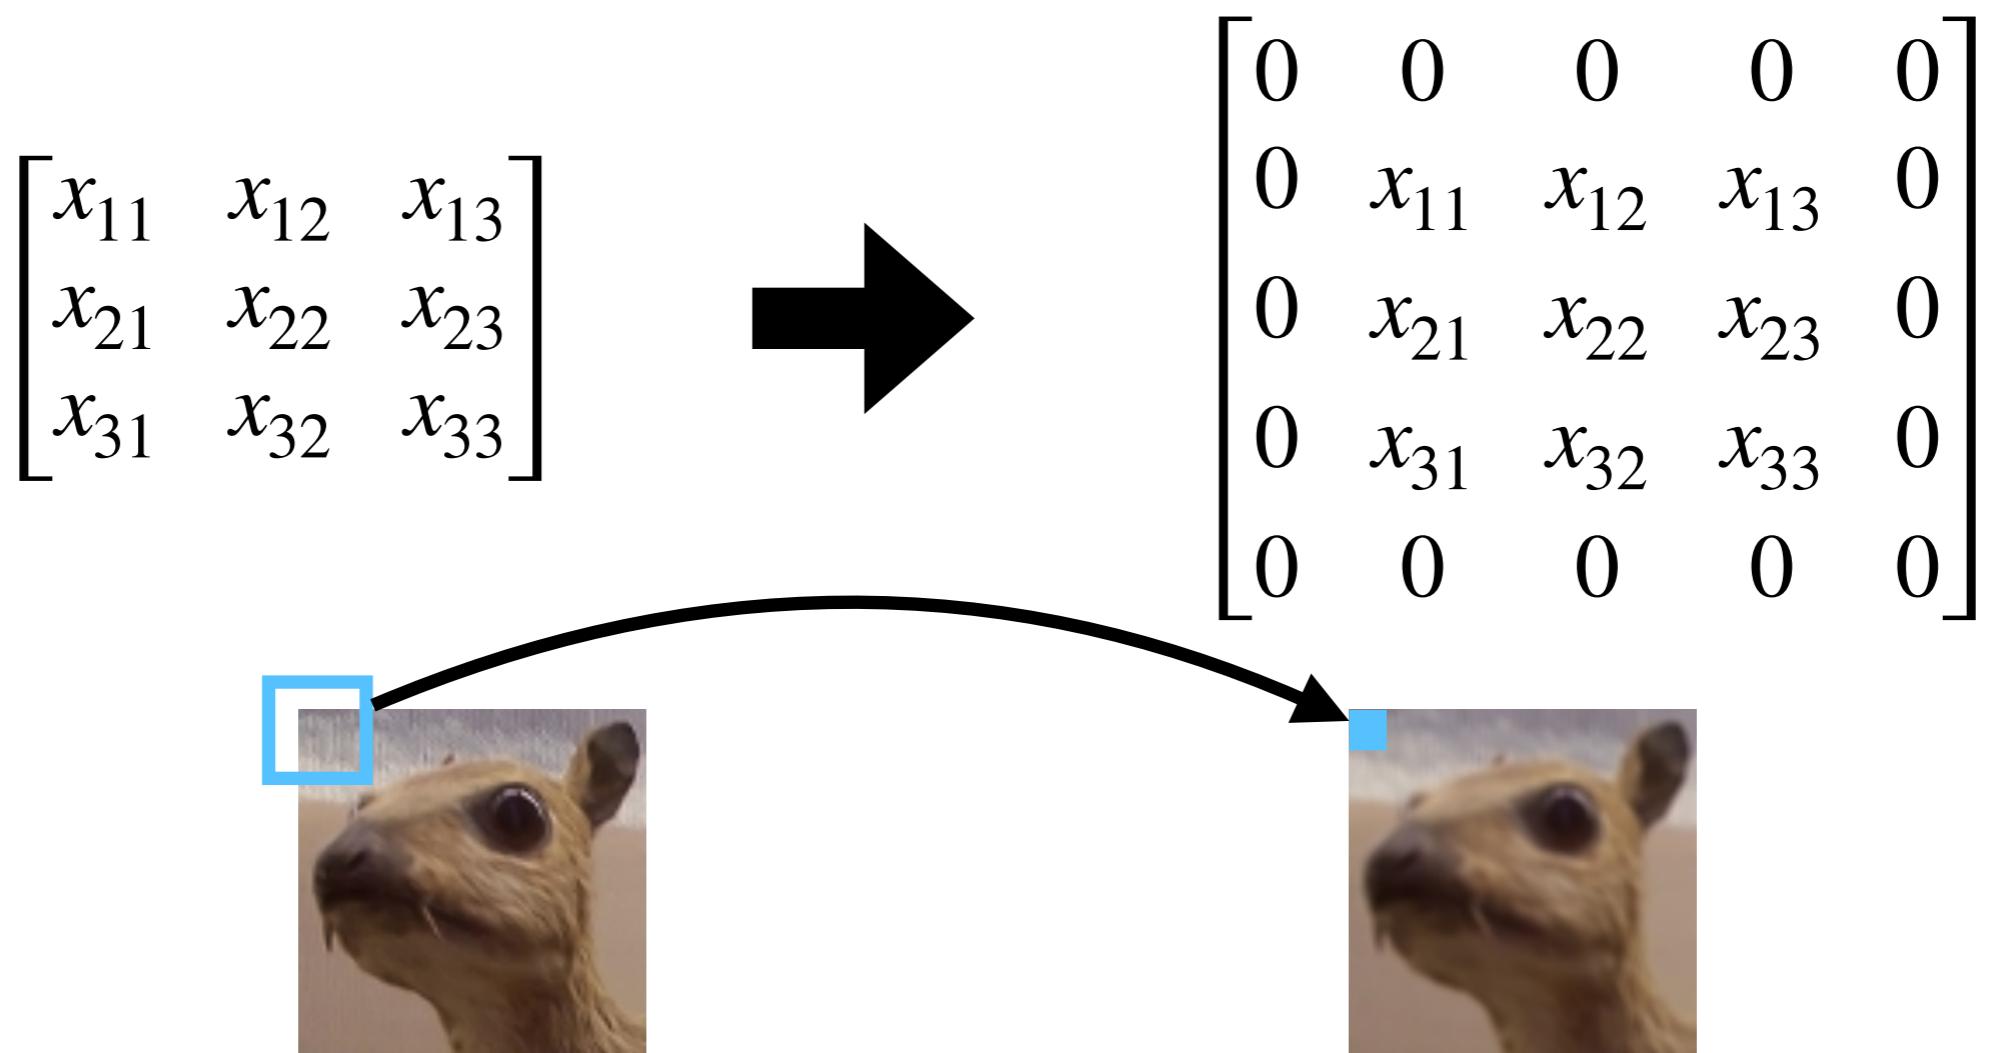
Problem

- The pixels in the corners of an image have minor contributions in the output



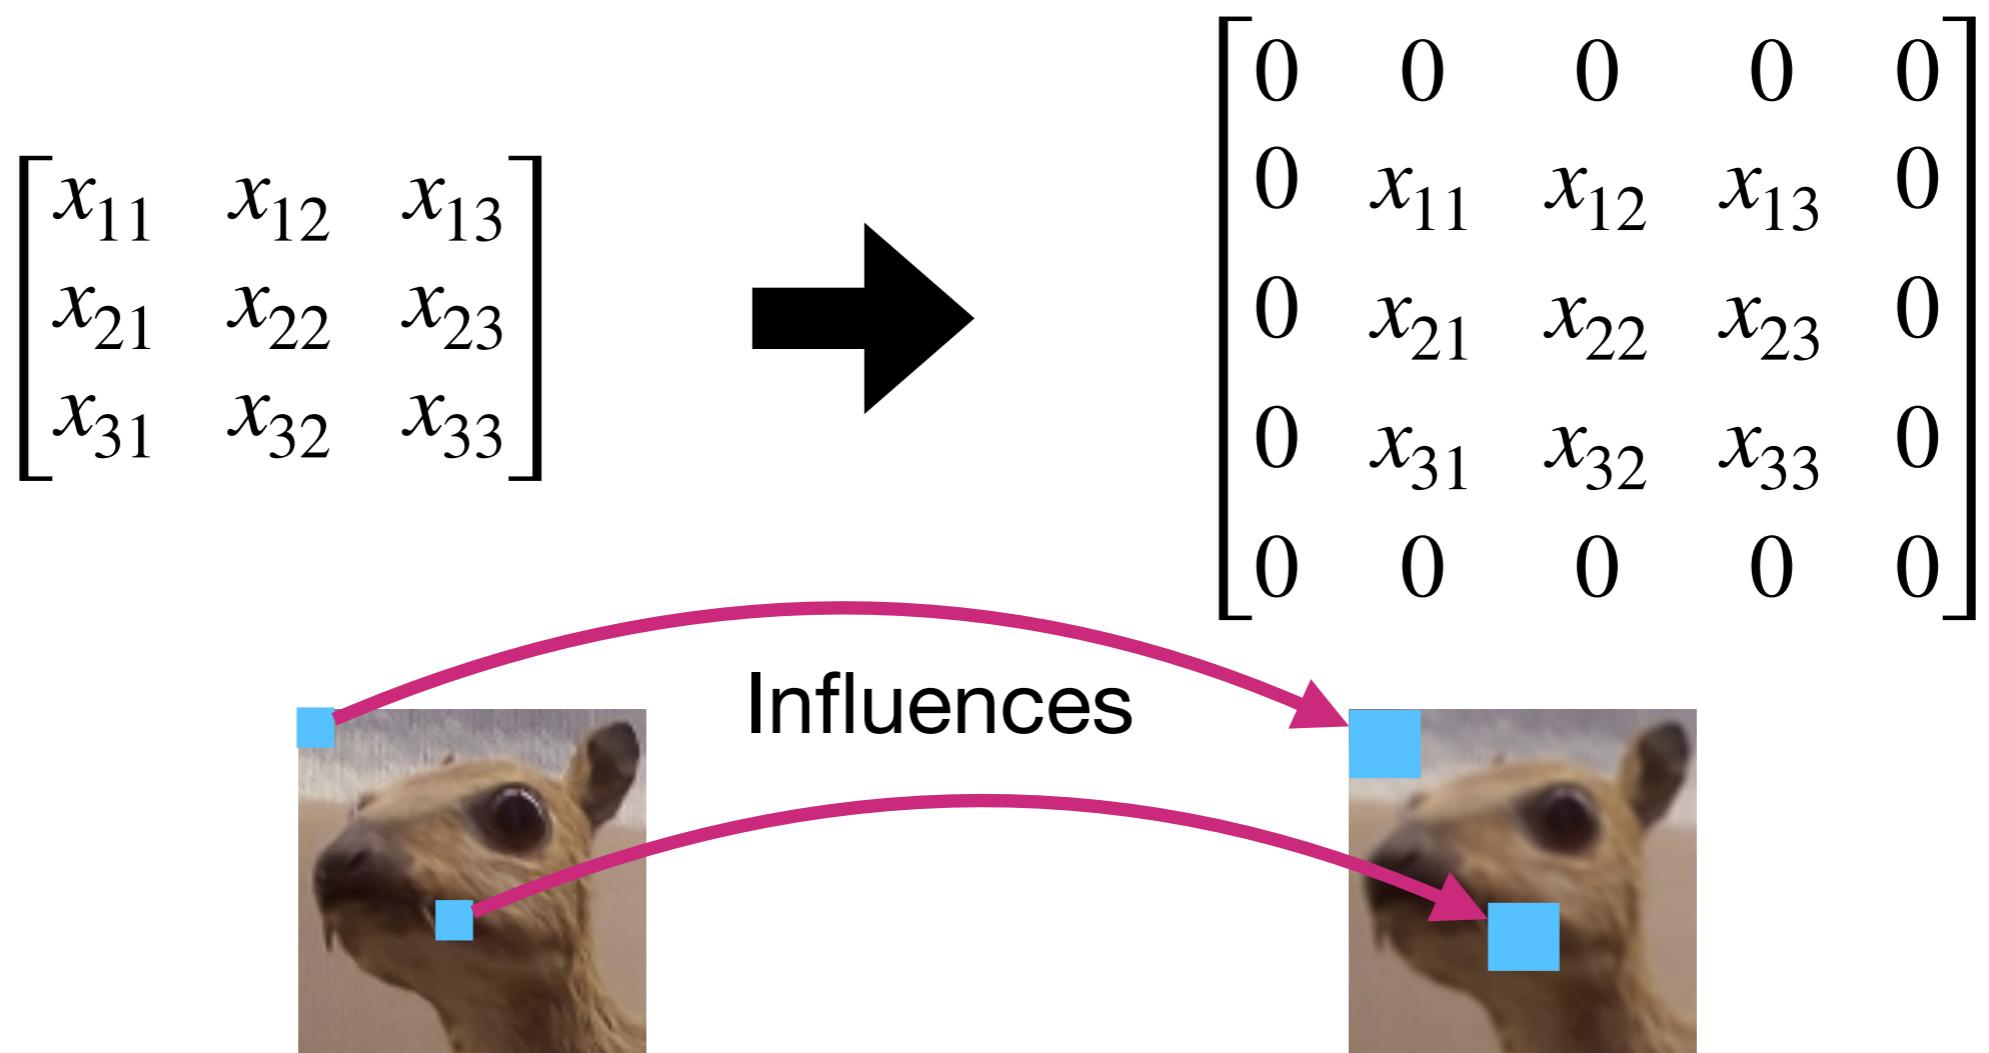
Padding

- Making matrices larger by padding with zeros

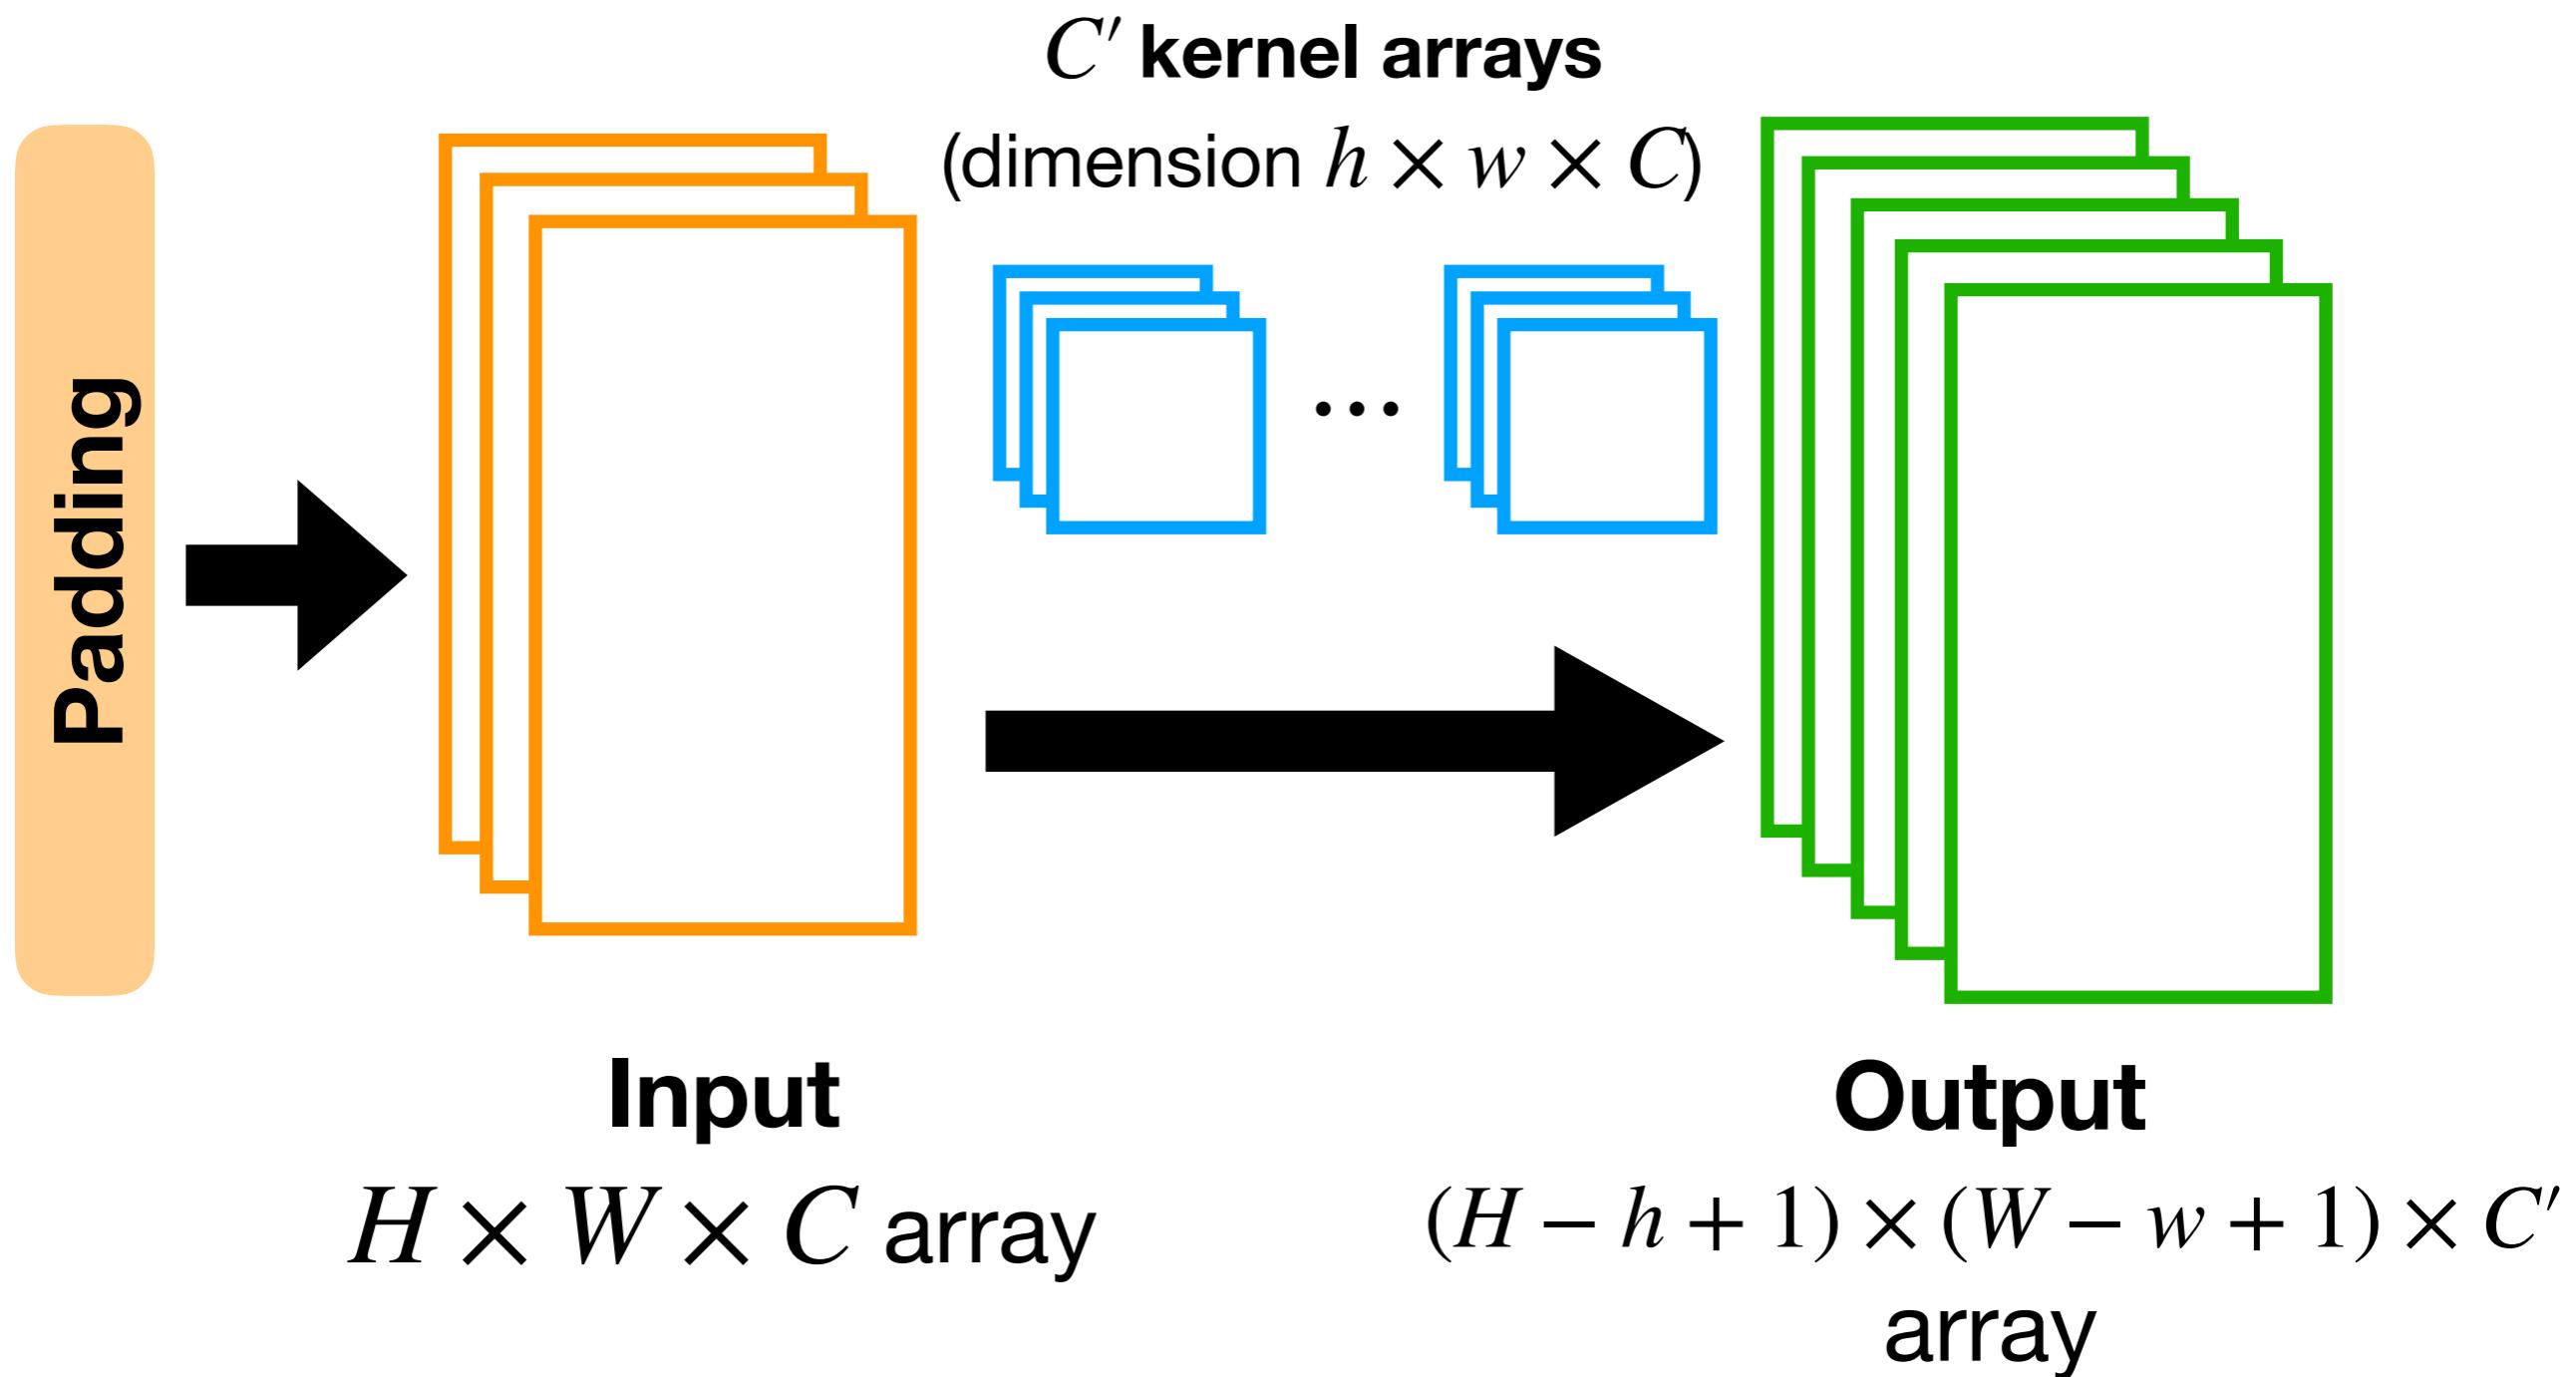


Padding

- Making matrices larger by padding with zeros



Convolutional layers



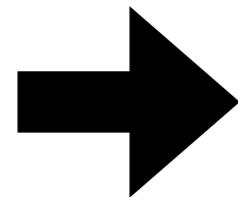
Convolution

Entries of kernel arrays are **learnable weights** optimized by training

C' kernel arrays

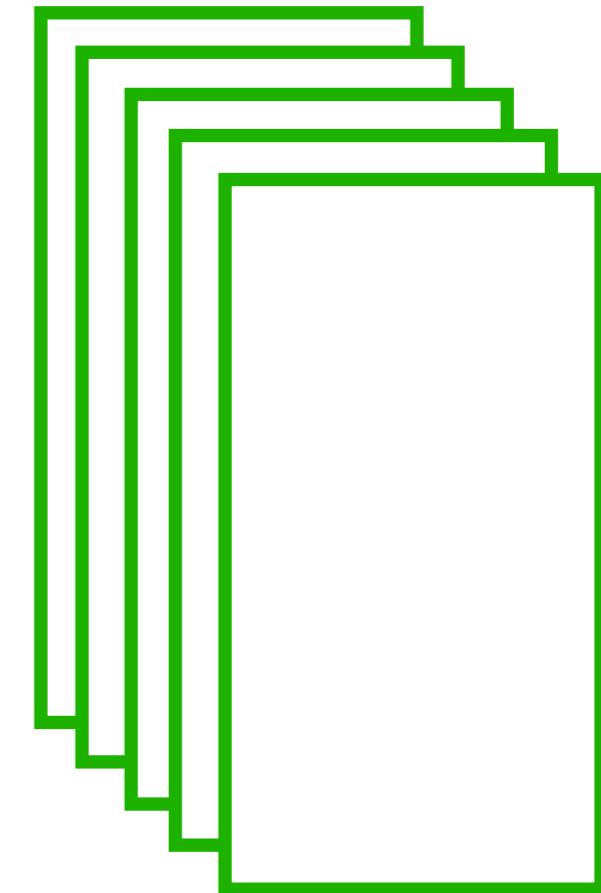
(dimension $h \times w \times C$)

Padding



Input

$H \times W \times C$ array



Output

$(H - h + 1) \times (W - w + 1) \times C'$ array

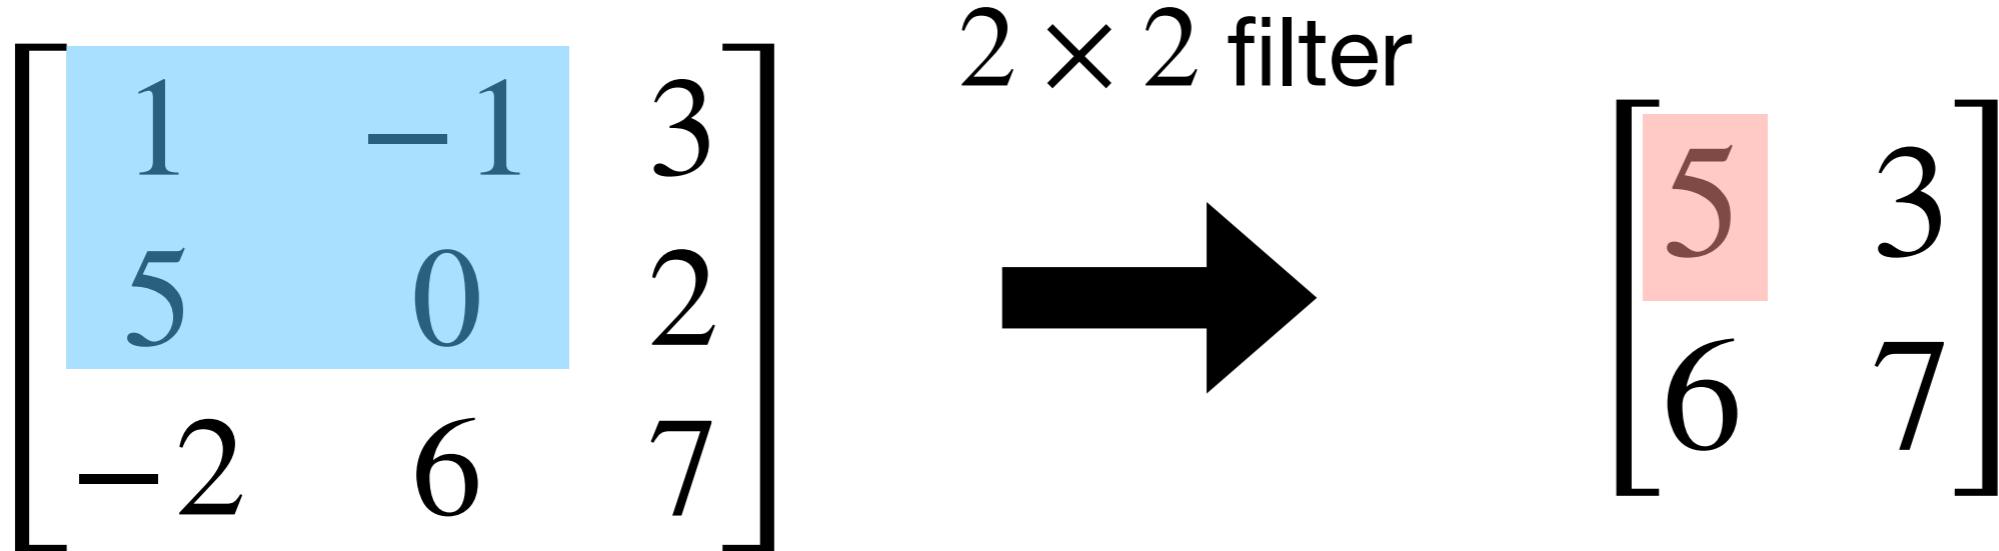
Pooling

- A technique to reduce the dimensions of features

$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix} \xrightarrow[2 \times 2 \text{ filter}]{} \begin{bmatrix} 5 & 3 \\ 6 & 7 \end{bmatrix}$$

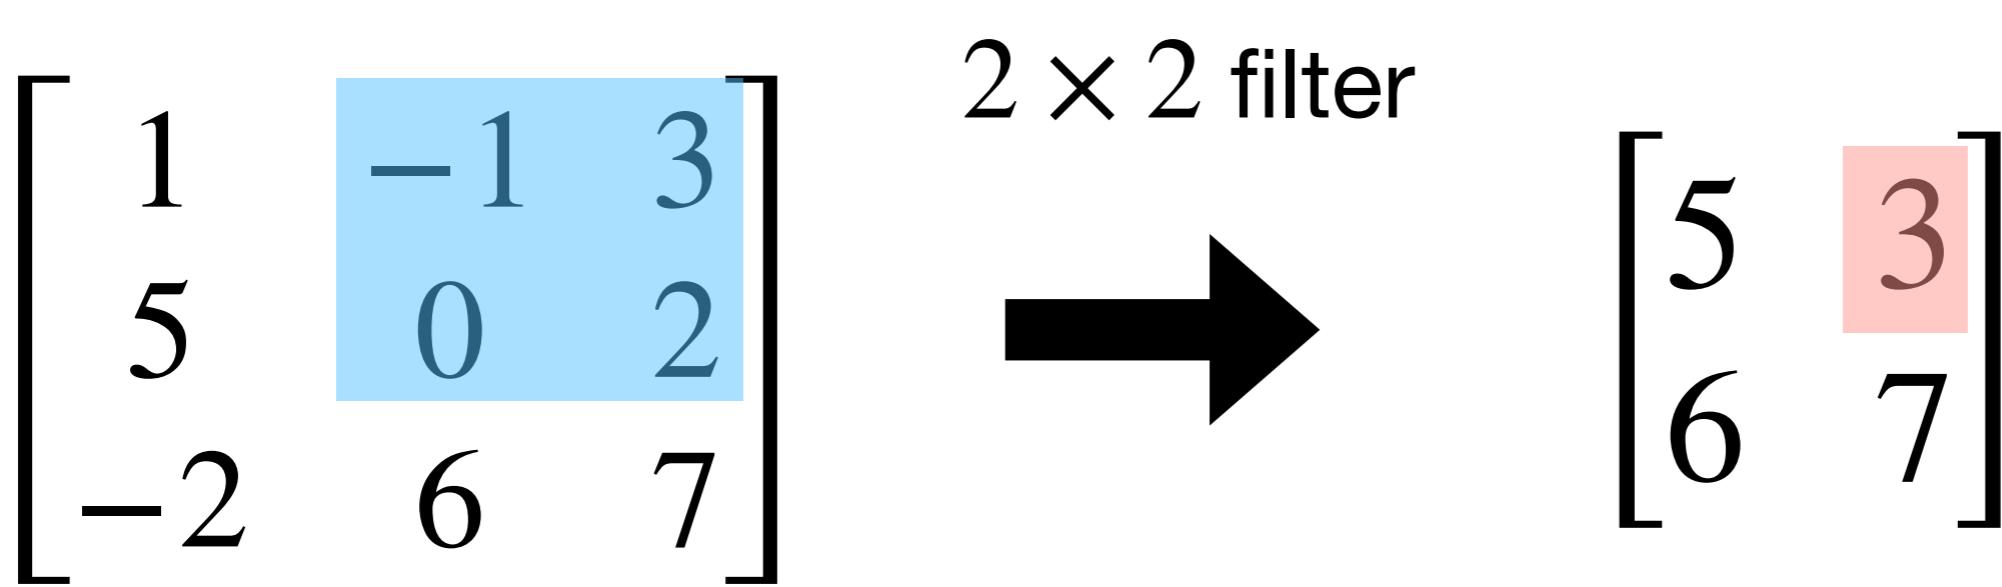
Pooling

- A technique to reduce the dimensions of features



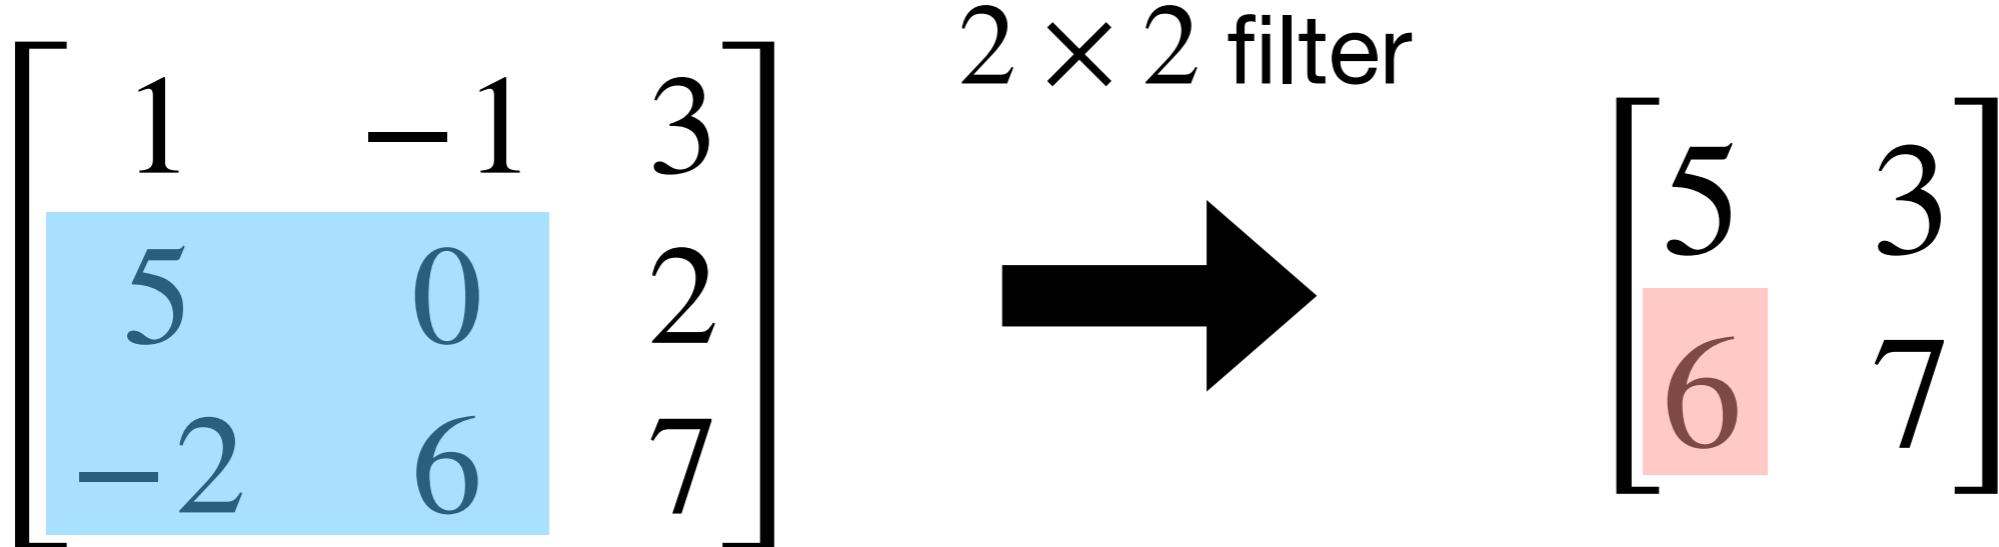
Pooling

- A technique to reduce the dimensions of features



Pooling

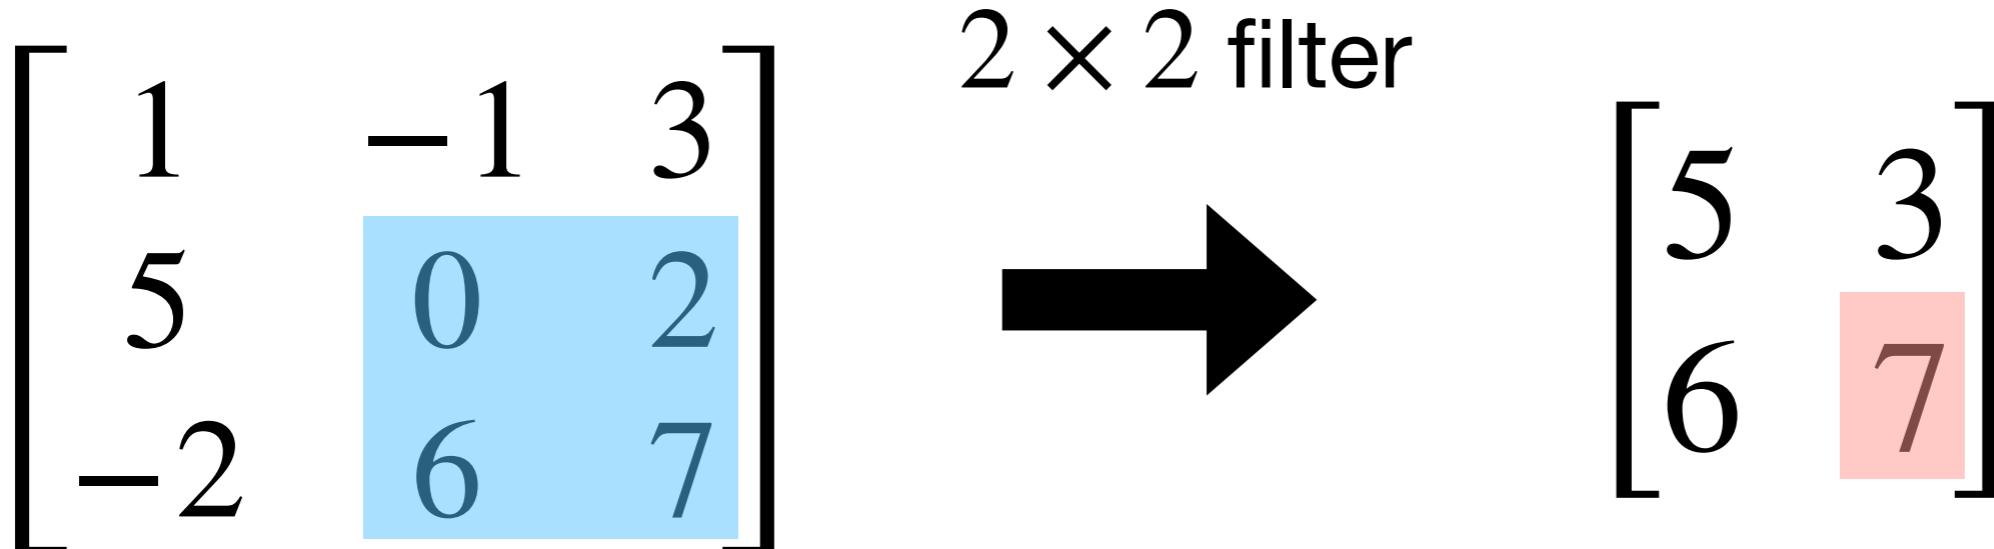
- A technique to reduce the dimensions of features



Pooling

- A technique to reduce the dimensions of features

Max pooling



Pooling

- A technique to reduce the dimensions of features

Average pooling

$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix} \xrightarrow[2 \times 2 \text{ filter}]{} \begin{bmatrix} 1.25 & 1 \\ 2.25 & 3.75 \end{bmatrix}$$

Pooling

- A technique to reduce the dimensions of features

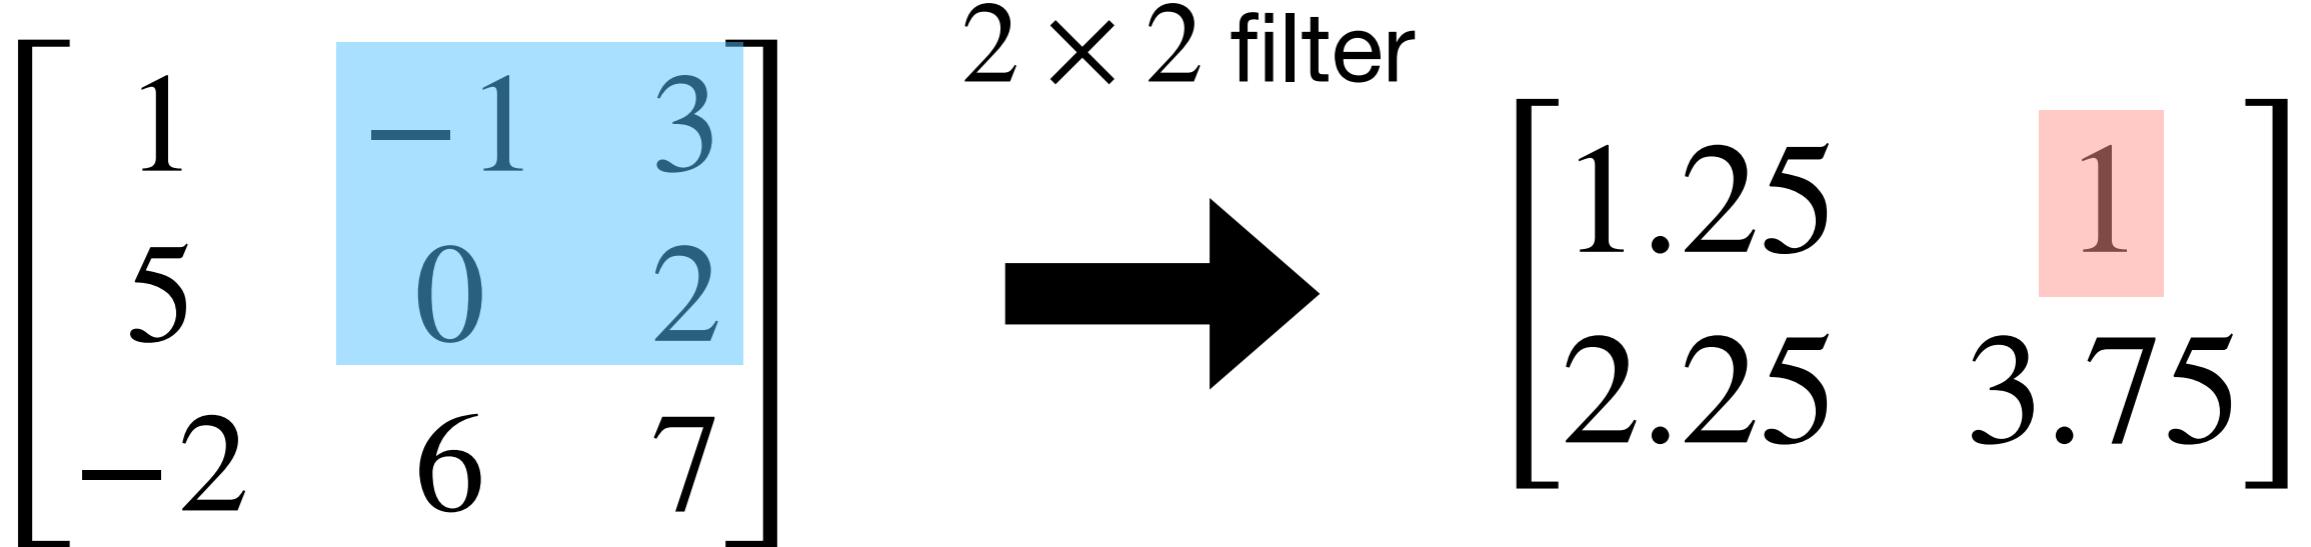
Average pooling

$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix} \xrightarrow[2 \times 2 \text{ filter}]{} \begin{bmatrix} 1.25 & 1 \\ 2.25 & 3.75 \end{bmatrix}$$

Pooling

- A technique to reduce the dimensions of features

Average pooling



Pooling

- A technique to reduce the dimensions of features

Average pooling

$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix} \xrightarrow[2 \times 2 \text{ filter}]{} \begin{bmatrix} 1.25 & 1 \\ 2.25 & 3.75 \end{bmatrix}$$

Pooling

- A technique to reduce the dimensions of features

Average pooling

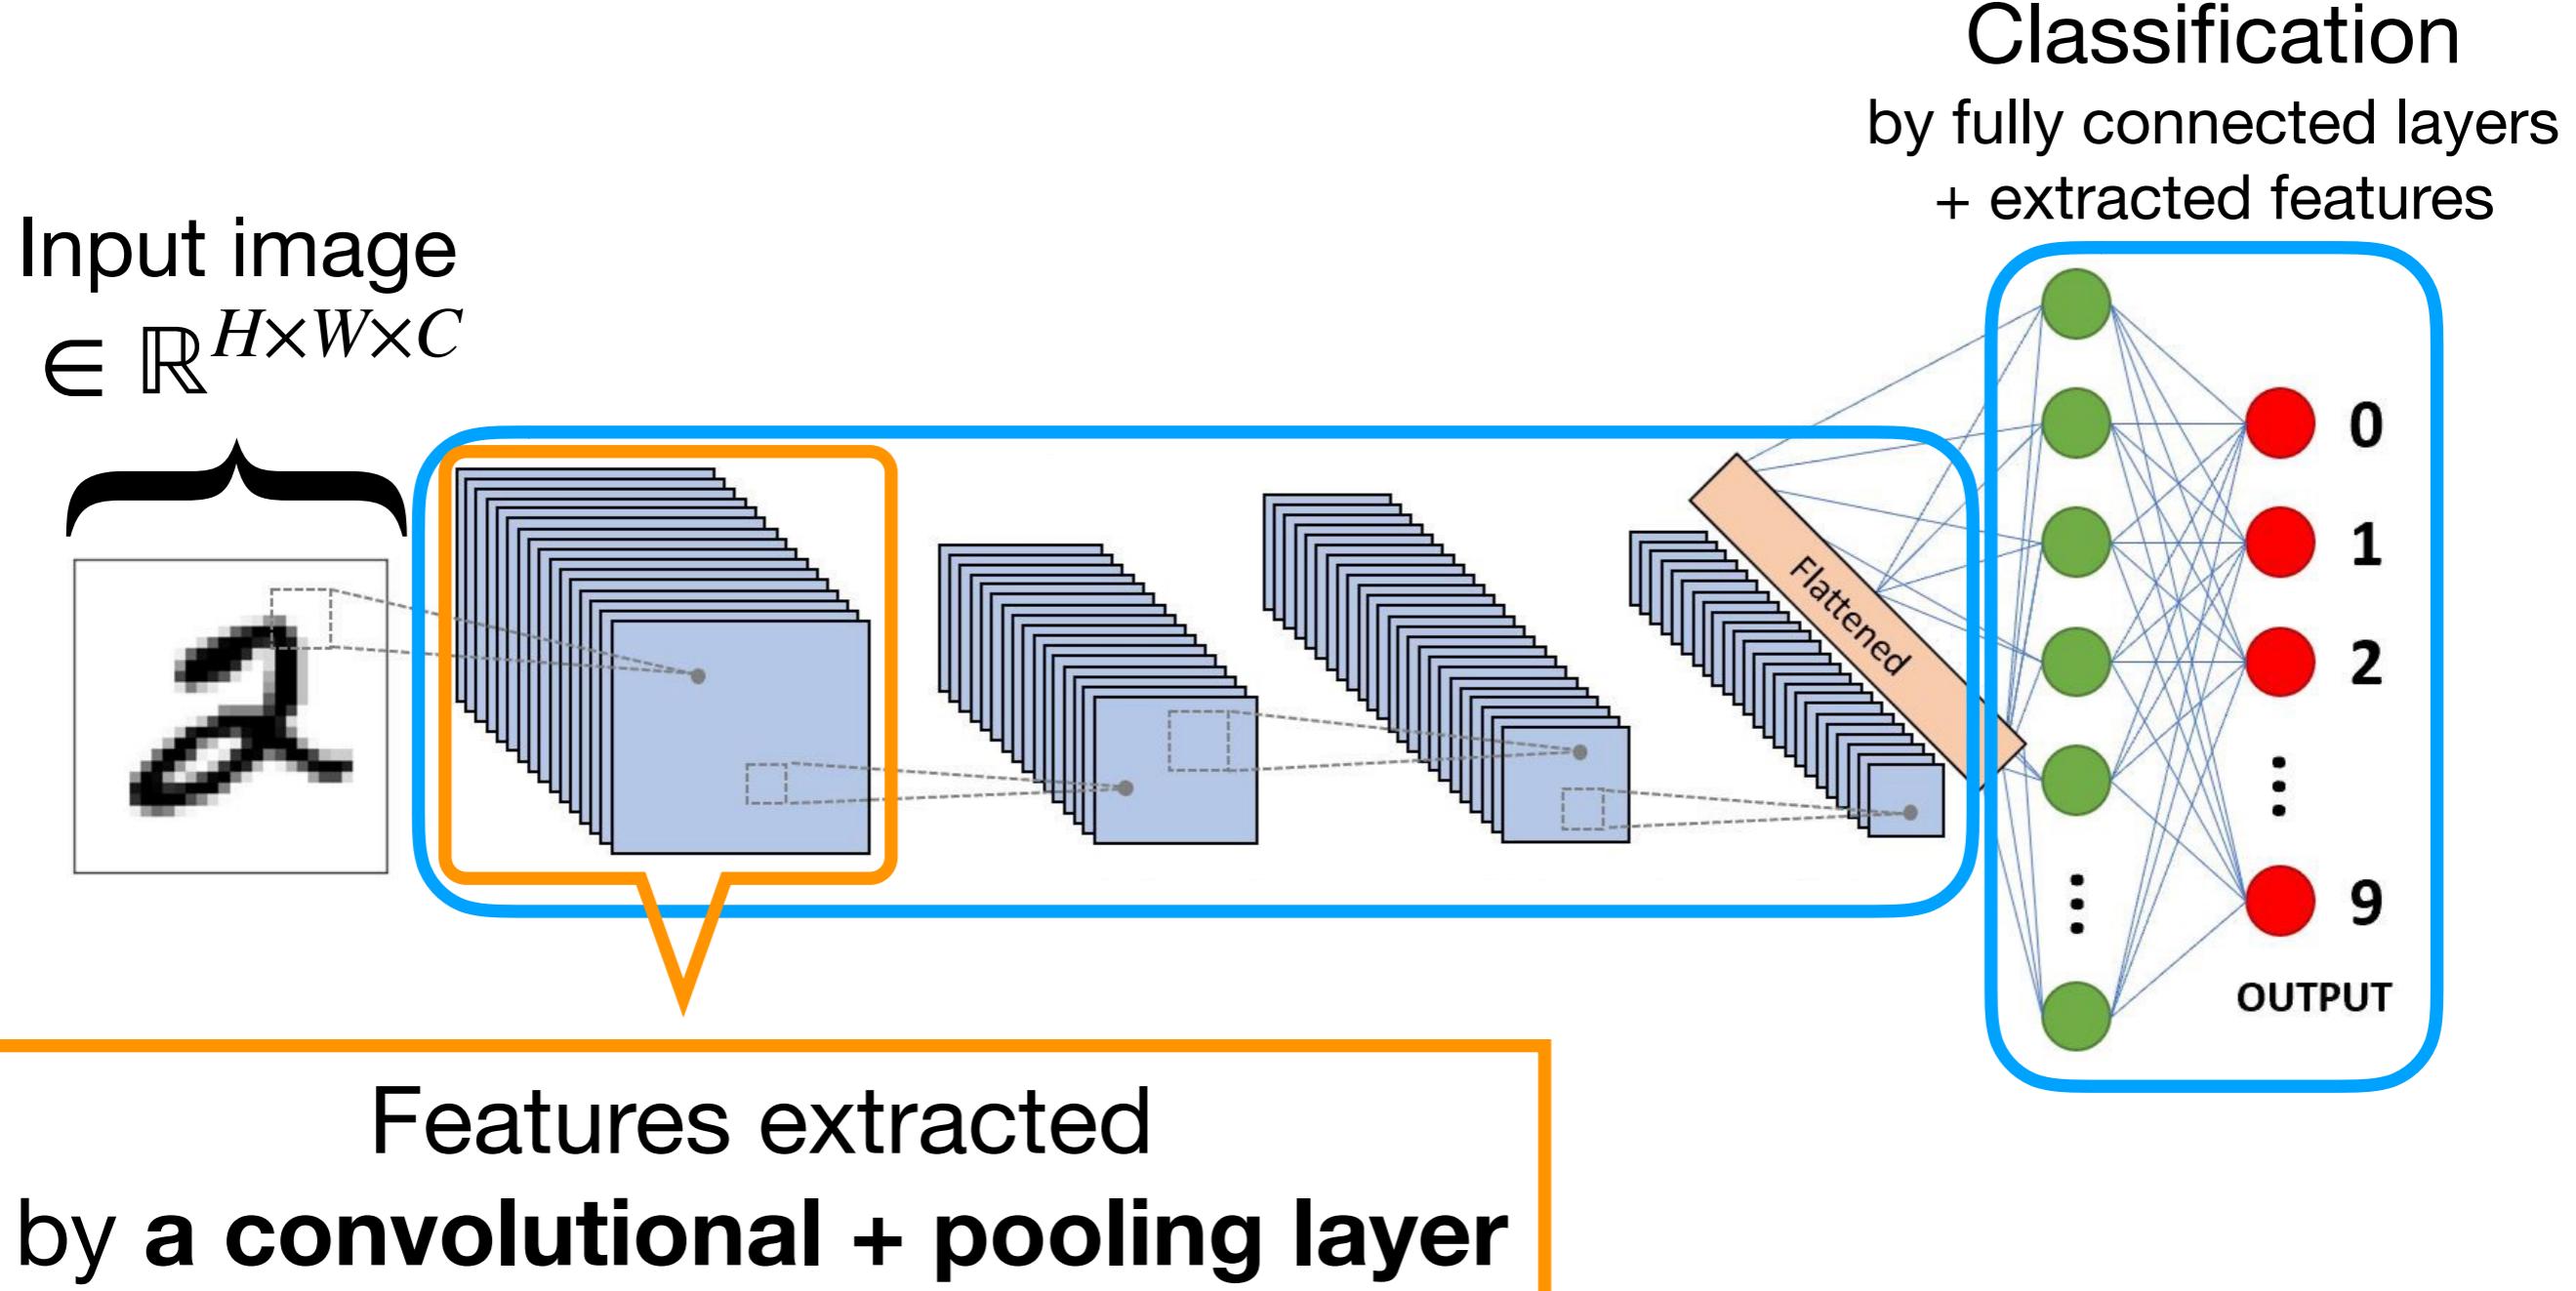
$$\begin{bmatrix} 1 & -1 & 3 \\ 5 & 0 & 2 \\ -2 & 6 & 7 \end{bmatrix} \xrightarrow[2 \times 2 \text{ filter}]{} \begin{bmatrix} 1.25 & 1 \\ 2.25 & 3.75 \end{bmatrix}$$

Why pooling?

Pooling layers contribute to:

- Speed up
- Simplifying models
 - Simpler models are more unlikely to be overfitted

Convolution neural networks



The image from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Other topics

- Dropout
- Residual blocks
- Batch normalization

This lecture

Introducing two special types of NNs

- Convolutional neural networks (CNNs)
 - Specialized in image processing
 - Involving layers to extract features from images
- **Recurrent neural networks (RNNs)**
 - Specialized in sequential data
 - Relaxing the vanishing gradient problem

Task examples with sequential data

- Natural language processing
 - Sentences are sequences of words
- Audio signal processing
 - Audio waveform is a sequence of signals
(in a digital format)
- Anomaly detection

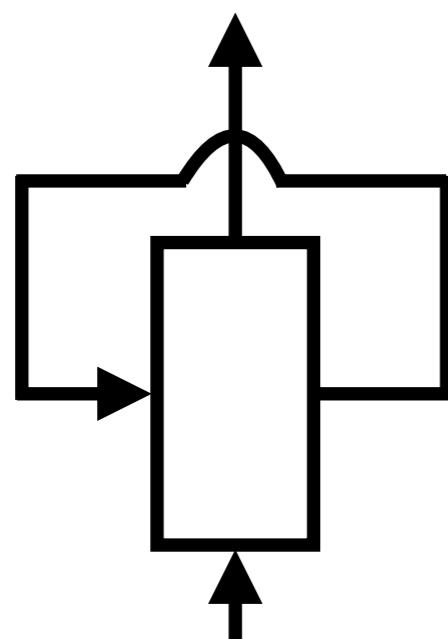
Challenge

- Sequential data have variable lengths
- Feedforward NNs cannot be used because they suppose the size of data to be fixed

Recurrent neural networks

- ***Expandable*** according to input data,
involving ***cyclic*** modules

Output features

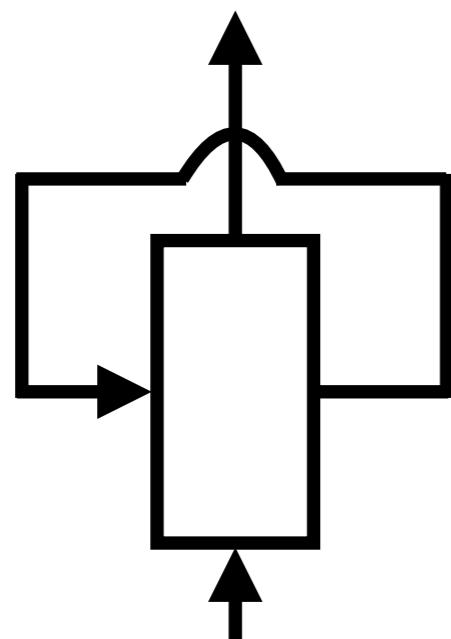


Input features

Recurrent neural networks

- ***Expandable*** according to input data, involving ***cyclic*** modules

Output features



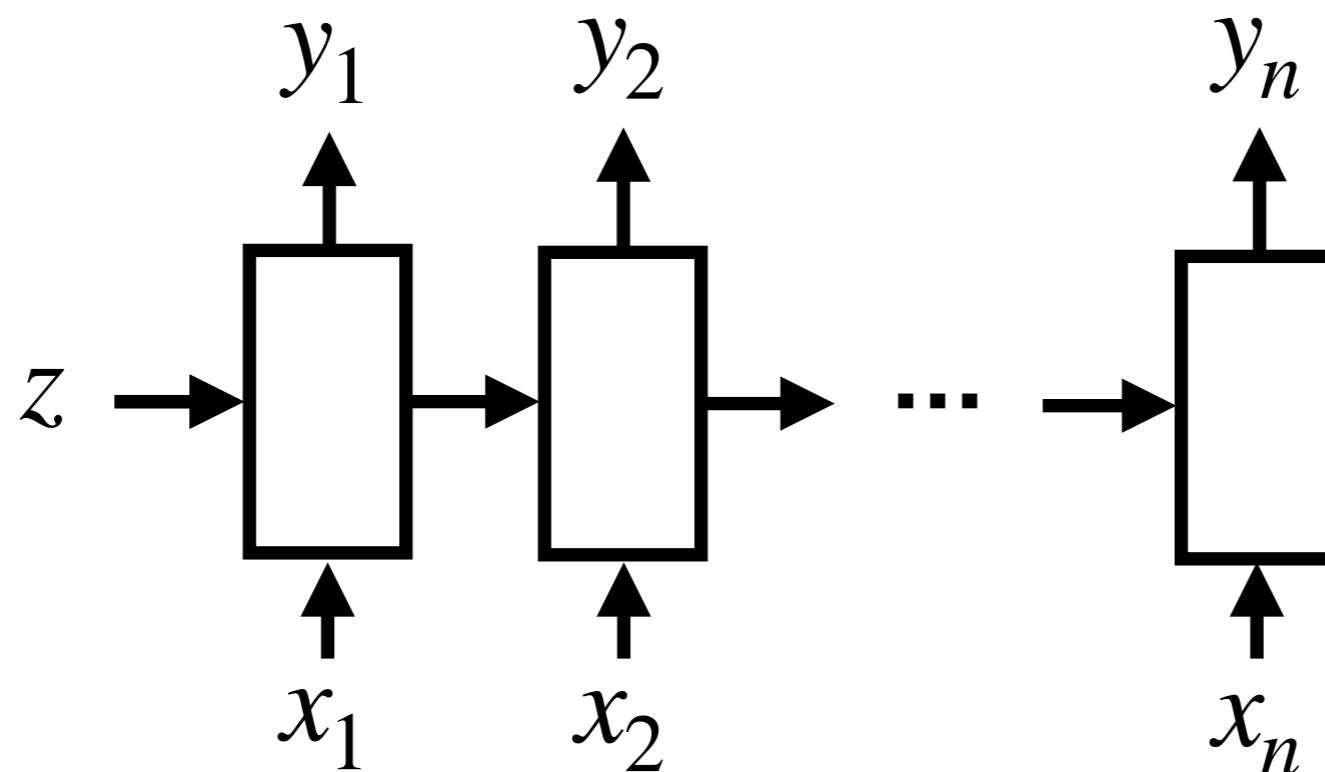
Input features

State features,
containing information propagated
from the current input to the next

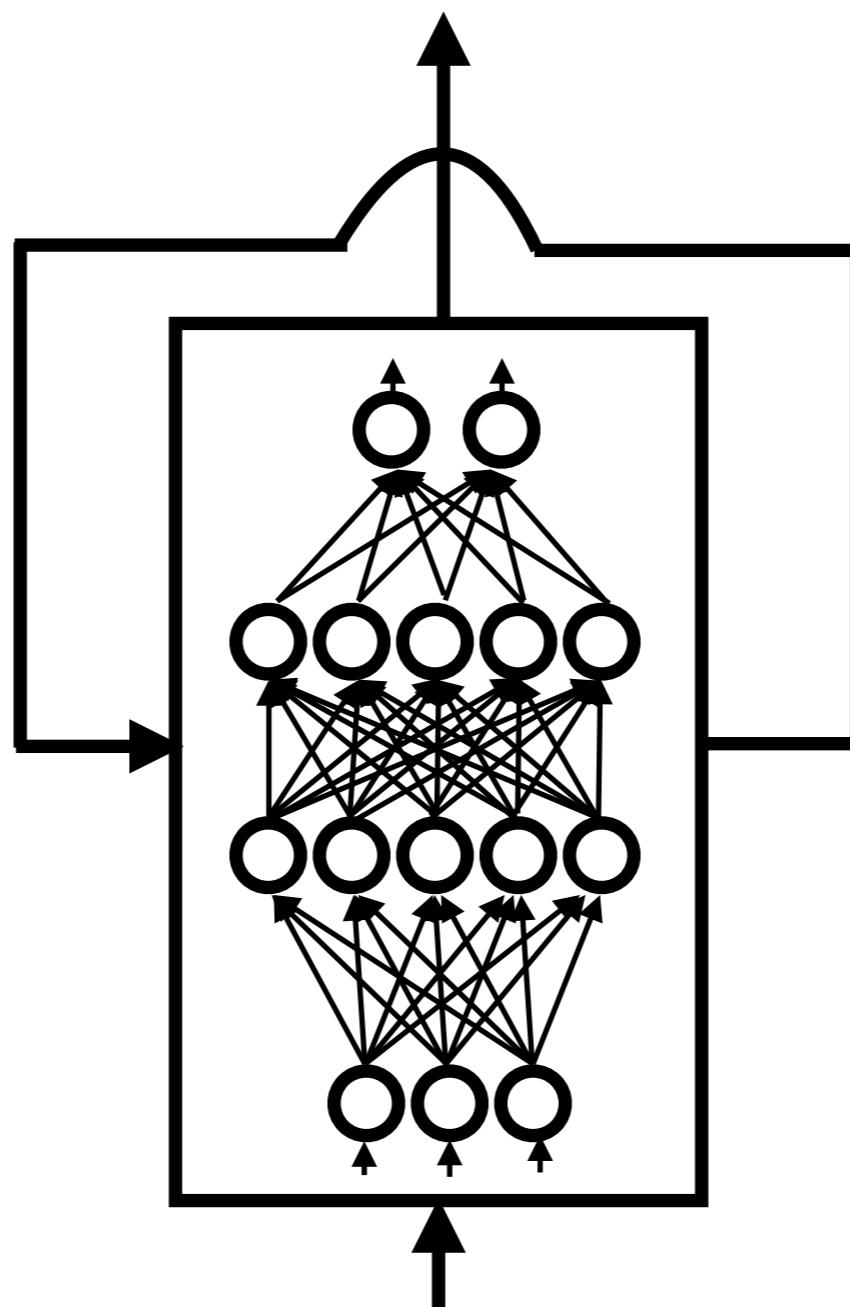
Recurrent neural networks

Given:

- Initial state feature $z \in \mathbb{R}^l$
- Input sequence $x_1, \dots, x_n \in \mathbb{R}^m$

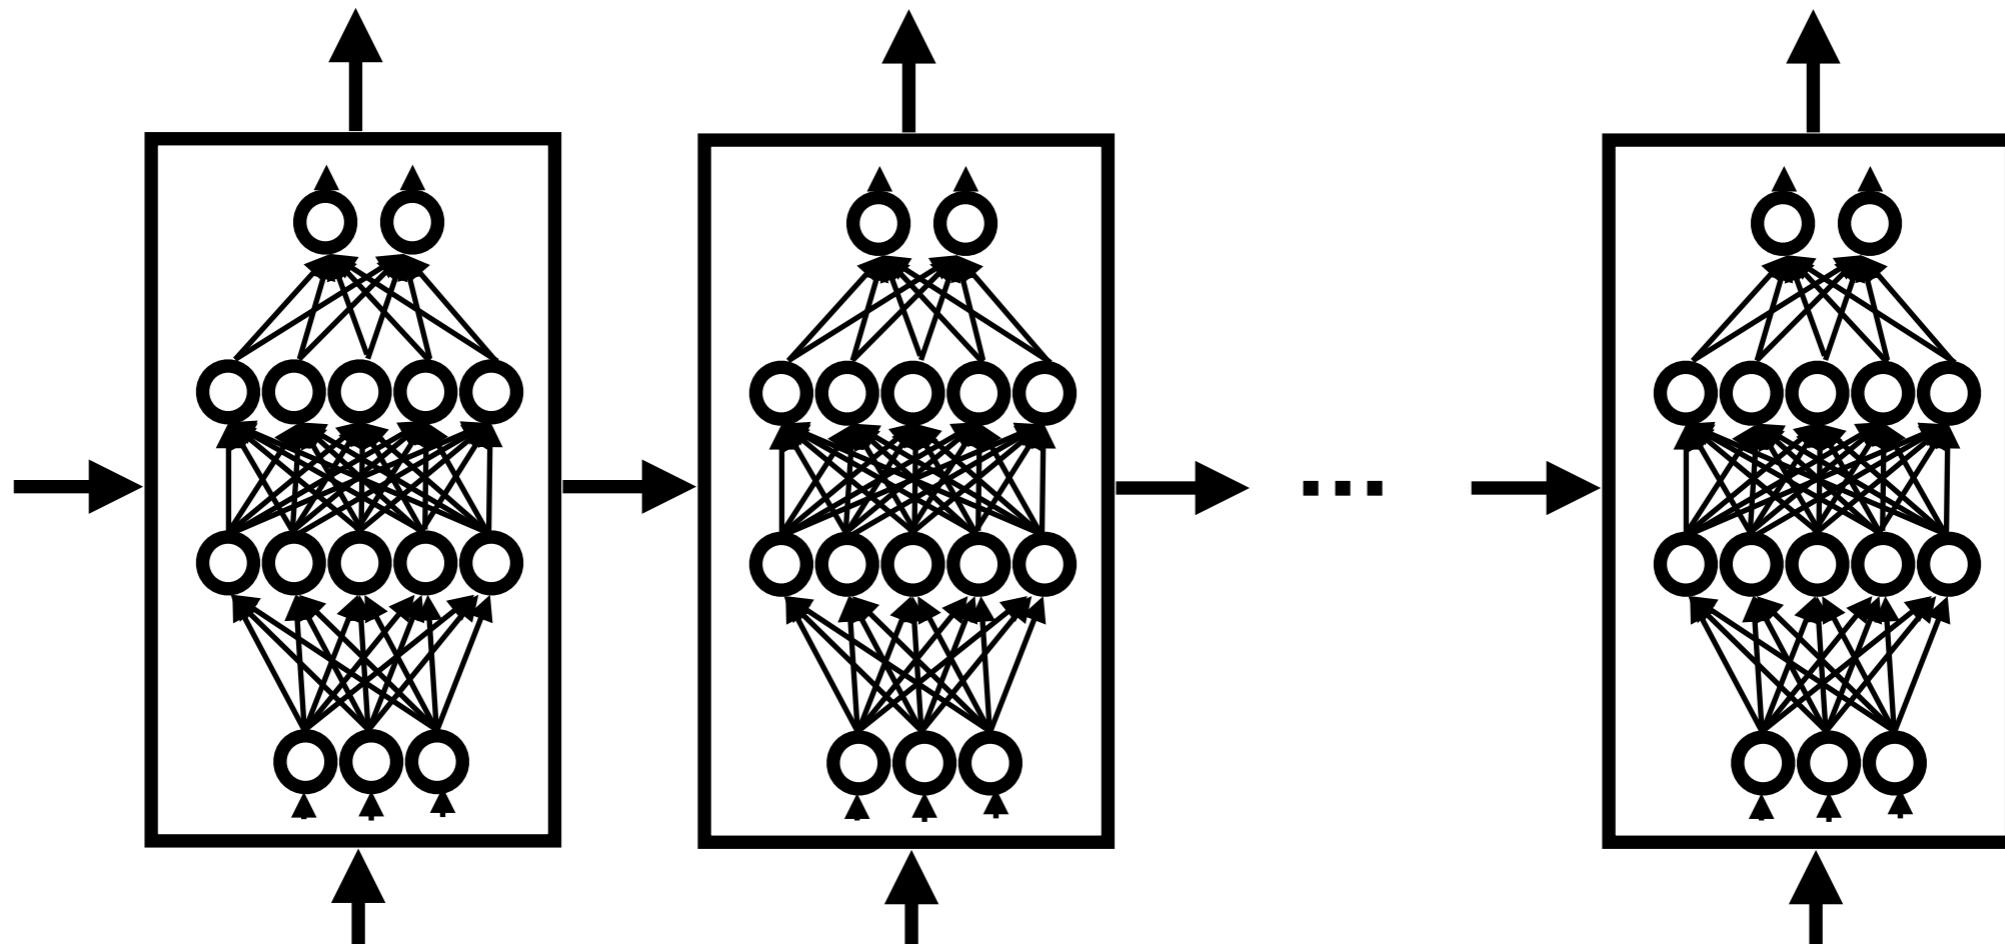


A simple implementation

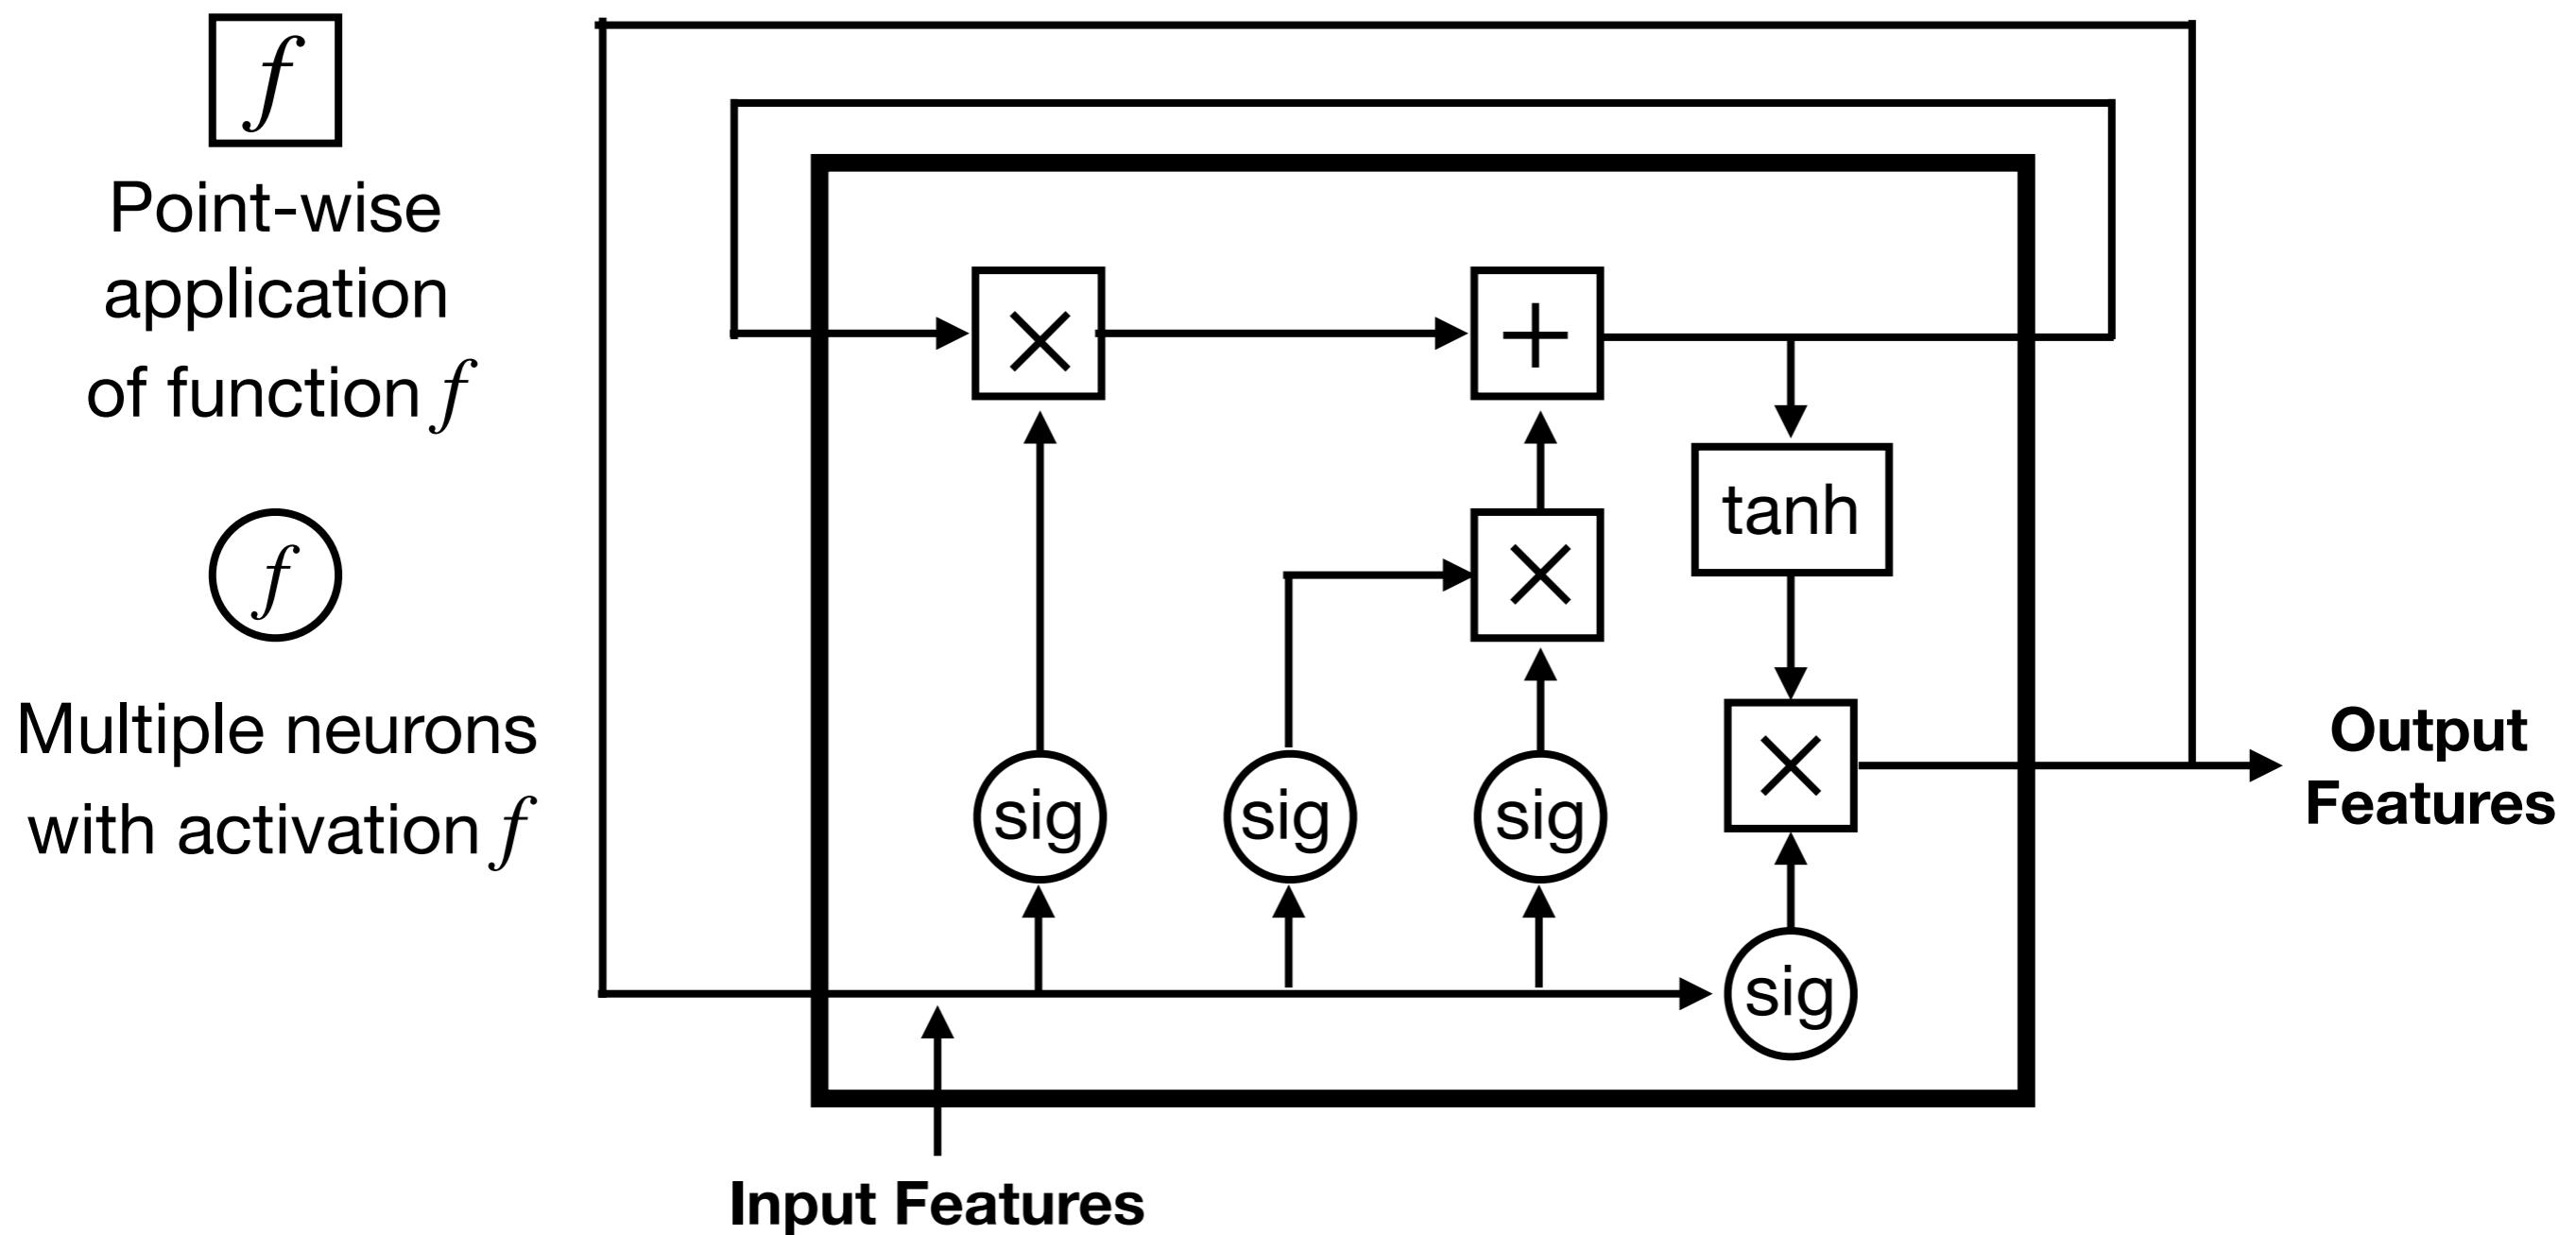


Problem

- Expanded NNs are very deep, so they suffer from the vanishing gradient problem



Long-short term memories (LSTMs)



Notice

- Next week is the turn of programming!
- Exercises will be assinged on Course N@vi